# WattDepot: An open source software ecosystem for enterprise-scale energy data collection, storage, analysis, and visualization

Robert S. Brewer
Philip M. Johnson
*Collaborative Software Development Laboratory*
*Department of Information and Computer Sciences*
*University of Hawai'i at Mānoa*
*Honolulu, HI 96822*
*rbrewer@lava.net, johnson@hawaii.edu*

*Abstract*—**WattDepot is an open source, Internet-based, service-oriented framework for collection, storage, analysis, and visualization of energy data. WattDepot differs from other energy management solutions in one or more of the following ways: it is not tied to any specific metering technology; it provides high-level support for meter aggregation and data interpolation; it supports carbon intensity analysis; it is architecturally decoupled from the underlying storage technology; it supports both hosted and local energy services; it can provide near-real time data collection and feedback; and the software is open source and freely available. In this paper, we introduce the framework, provide examples of its use, and discuss its application to research and understanding of the Smart Grid.**

## I. INTRODUCTION

Recent interest in the Smart Grid has produced a wide variety of proposed "open" standards and technologies for energy data collection, storage, and analysis. Some prominent approaches include cloud-based services like Google Power-Meter [1] and Microsoft Hohm [2], "interoperable" protocols and networking approaches like Smart Energy 2.0 [3] and GridRouter Ecosystem [4], smart grid standards like IEEE SCC21 [5] and Oasis Blue [6], extensible energy management devices like Control4's EMS 100 [7], and home energy meters like the TED 5000 [8]. There are also "open source" hardware platforms, such as OSHAN [9], and open source software solutions for phasor data such as OpenPDC [10].

For our research on behavioral change among energy consumers in a building or home environment, we need a way to:

1) collect energy data from a variety of meters with the possibility of near-real time (5-10 second) feedback
2) store the results in an Internet-accessible repository
3) perform basic analyses on the raw data including aggregation and interpolation
4) visualize the raw and processed data in a variety of ways, including tables, gauges, trend lines, geographic maps, heat maps, and so forth.

Unfortunately, we find that none of the current "open" technologies satisfy these seemingly simple requirements. Either the technology is designed for a specific brand of meter (such as the EMS 100 or TED 5000), or it does not support near-real time feedback (such as Google PowerMeter and Microsoft Hohm), or it focuses on utility issues (OpenPDC), or it focusses on wire-level or hardware concerns (OSHAN, GridRouter Ecosystem).

Rather than implement a special purpose solution for our research, we leveraged our prior experience in open source software engineering measurement systems to design and implement an open source, extensible, service-oriented framework for energy data collection, storage, analysis, and visualization. Our framework, called WattDepot, consists of three kinds of services:

- WattDepot *sensors*, each customized for a particular brand of energy meter. A sensor requests data from a meter according to the meter's protocol, then sends it to a WattDepot repository for storage.
- WattDepot *servers*, which implement a REST [11] API for accepting energy data sent from sensors and providing this sensor data (or analyses based upon the data) to WattDepot clients.
- WattDepot *clients*, which request data from WattDepot servers and either display the data or analyses directly to users or provide the data to higher level energy services.

Fig.1 illustrates the architecture of the system.

Consider the following application example: a University wants to implement a dorm energy competition in which energy meters will be installed on each floor of the dorm and residents of each floor will compete against each other to see who can reduce their energy the most. Prior to WattDepot, one could either buy a commercial, turnkey solution such as the Building Dashboard from Lucid Design Group and be limited to the set of meters and analyses they support, or else start from scratch and build all of the energy data collection, analysis, and visualization services. WattDepot provides an alternative approach, in which collection, storage, and analysis capabilities are freely available and amenable to integration
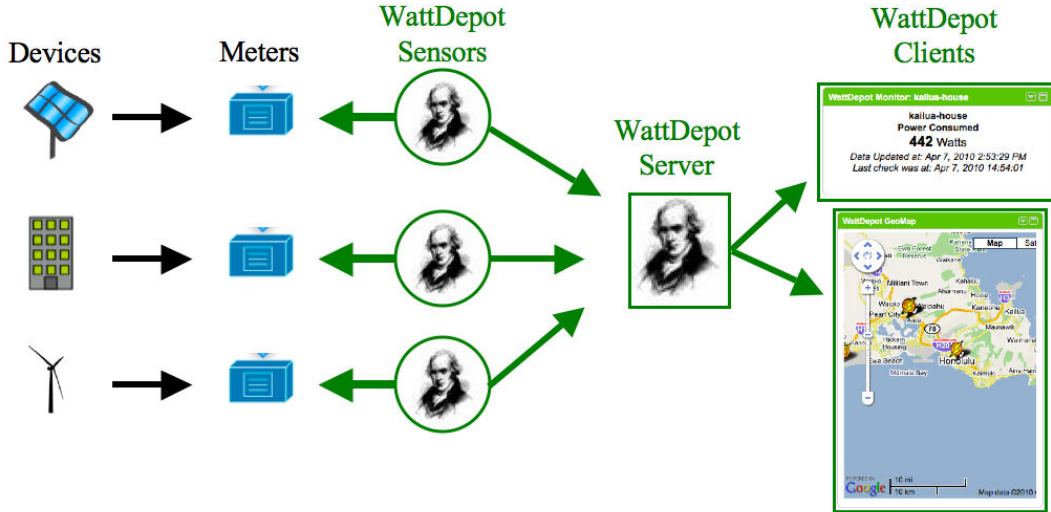
Fig. 1. *Architecture of WattDepot: sensors obtain data from meters on consumption or generation, transmit their findings to a server, which is queried by clients to present visualizations or analyses.*

within a web application developed by the user. For those with some software development capability, this system can be a lower cost solution and/or support customization beyond the ability of current proprietary solutions.

WattDepot services are designed at the "enterprise" scale as opposed to the "home" or "utility" scales. By home scale, we mean systems like the TED 5000, which are designed to collect and analyze data about a single residence, typically through a simple "dashboard" interface that runs on the user's local area network. WattDepot is designed to collect, store, and analyze data from a wider range of locations using the Internet as the transmission mechanism, and our performance analysis tests indicate that it can easily process several thousand sensor data storage requests per minute given adequate network bandwidth. On the other hand, by utility scale, we mean systems that can collect and store grid-level data comprising tens or hundreds of thousands of residences, with revenue-grade measurement, hardened storage, and security mechanisms. WattDepot is not intended for these more rigorous utility-level requirements. Our goal for WattDepot is to provide a mechanism to accelerate innovative research and development in energy analysis and visualization by providing an extensible architecture and well documented, open source software platform.

The next section of this paper introduces the basic components of WattDepot. Following this we describe some of our initial experiences with the framework and conclude with our future directions.

## II. WATTDEPOT SERVICES

### A. Sensors

Electricity can be generated and/or consumed by a wide variety of devices and monitored by a wide variety of metering technologies. Many energy management software systems are tied to a particular energy device and/or meter; indeed, they are often marketed as an accessory to the hardware. For example, a solar panel company might provide software for storing data about the power generated by the panel and simple visualizations of the power generated over time.

The design of WattDepot attempts to make the software as independent from the energy hardware devices as possible. In WattDepot, it is the device that is the accessory to the software. To achieve this independence, the WattDepot architecture includes "sensors", or small software processes that query any given energy device according to its native protocol, collect standard information, and then send it to a WattDepot server using the Internet and the RESTful WattDepot API over HTTP. This design means that new energy devices can be integrated easily into the WattDepot software ecosystem just by writing the sensor interface. The use of a RESTful HTTP protocol for communication with the WattDepot server means that the sensor can be implemented in any programming language.

We have implemented WattDepot sensors for the TED 5000 home energy device, the Veris power meter collected through the Building Manager Online system, and for the Acuvim energy meter. This latter sensor is particularly interesting because it communicates with the meter through the standard ModBus/TCP protocol. The sensor is designed in a modular way so that it is easy to customize to support other meters that communicate using ModBus/TCP.

Sensors transmit data to a WattDepot server using a standardized, published protocol, as discussed next.

### B. Servers

A WattDepot server accepts raw energy data from devices (via sensors) and makes this data (or analyses based upon it) available to clients. The WattDepot server implements a variety of design decisions intended to improve its generality, reusability, and extensibility, including: (1) a RESTful API, (2)

a pluggable back-end database, (3) aggregation via "virtual" sources, (4) data interpolation, and (5) multiple representations.

*RESTful API.* WattDepot conforms to modern web service design best practices by providing a RESTful API. REST (REpresentational State Transfer) [11] is a specification paradigm, which, when applied to web services, generally results in more easily usable and extensible communication than alternatives such as SOAP. The details of RESTful design are beyond the scope of this paper. Some of the implications include URLs that can serve as unique identifiers for energy data and the use of HTTP methods such as PUT, POST, GET, and DELETE to add, amend, retrieve, and delete energy data. The WattDepot API [12] provides more details and a full specification of the supported operations.

*Pluggable back-end database.* The WattDepot server implements an abstraction layer that enables the server to be built with a variety of different persistence mechanisms. By default, WattDepot uses the Apache Derby relational database, which is a high performance, embedded database written in Java. However, WattDepot can be ported to other relational or non-relational database systems by implementing an interface and setting some run-time configuration parameters.

*Aggregation via virtual sources.* The data from a single physical meter is generally represented in WattDepot as a "Source". Each WattDepot Source can indicate the power generated or consumed by that device at any moment in time, the energy generated or consumed by that device over a period of time, the carbon intensity associated with that device, and other features. In addition to this one-to-one correspondence, WattDepot also supports the definition of "virtual" Sources, which are Sources defined as the aggregation of other Sources. For example, a floor on a building might have two meters collecting energy consumption data for the two sections of the floor. WattDepot allows users to define a virtual Source representing the aggregation of the data stream from the two meters. This virtual Source thus represents the total energy consumption for the floor. Virtual Sources can be arranged in a hierarchy, such that the virtual Sources for each floor in the building can themselves be contained in a virtual Source representing the entire building's energy consumption.

*Data interpolation.* One common initial roadblock to analyzing energy data collected from multiple sources is the "timestamp problem". For example, assume that energy consumption on a building floor is collected by two meters, and that energy data is collected from those meters becomes available approximately every 15 minutes, but the timestamps associated with the data for a given time period differ by a minute or two. Determining the aggregate energy consumption is no longer a simple matter of importing the two data sets into a spreadsheet and using a summation macro, because the timestamps from the two meters do not "match up". WattDepot addresses this problem by providing automatic interpolation. For example, assume a meter sent energy data at roughly 30 minute intervals: 11:23 AM, 11:56 AM, 12:25 PM, and 1:01 PM. You can request the energy consumed by this Source

between 12:00 PM and 1:00 PM, for example, and WattDepot will automatically interpolate the raw data values to provide an estimate for the interval of interest. Automatic linear interpolation enables virtual Sources to return reasonable values even when its constituent Sources send data at different times and with different frequencies. Fig.3 illustrates both virtual sources and data interpolation.

*Multiple representations.* One benefit of a RESTful architecture is the separation of a "resource" from its "representation". In WattDepot, this separation means that data and analyses for a given Source can be provided in different ways, depending upon the needs of the client. So far, WattDepot can provide data to clients using XML, JSON (the format supported by Google Visualizations), and CSV (comma-separated values), which is useful for importing WattDepot data into other tools for additional analysis.

## C. Clients

While sensors collect energy data, and the server stores, aggregates, and interpolates it, WattDepot clients extract the data for presentation to the user or input into tools for additional analysis. The WattDepot RESTful API is designed to enable a wide variety of clients to interact easily with WattDepot energy data.



Fig. 2. *The WattDepot monitor client displaying the latest sensor data for a home.*

Fig.2 and Fig.4 illustrate two WattDepot clients, both based upon the Google Gadget technology which enables users to easily create personalized "dashboards" containing a wide variety of information resources. Fig.2 shows a client that provides a real-time monitor. After installing the gadget, users configure it with the URL of a WattDepot server, the Source that they wish to monitor, the type of data to display, and an update interval (between 5 and 30 seconds). The gadget queries the WattDepot server according to the update interval and refreshes the gadget window with the most recent data.

The client in Fig.4 is a map-based perspective on Sources. Users can define a Source with its latitude and longitude coordinates, and this information can be used to display the set of Sources associated with a WattDepot server according to their location. After installing the gadget, users configure it with the URL to a WattDepot server. The gadget then displays icons on the map for each Source in the server. Clicking an icon triggers the gadget to retrieve the most recently received data by that Source and display it in a pop-up window.

Fig.3 shows a "Visualizer" web application we built as part of a developer toolset. Users select one or more sources,

Fig. 3. *The WattDepot Visualizer client displaying a virtual Source (upper trace) for a building floor and its two constituent physical Sources (in lower traces) corresponding to energy meters. All data values are interpolated.*
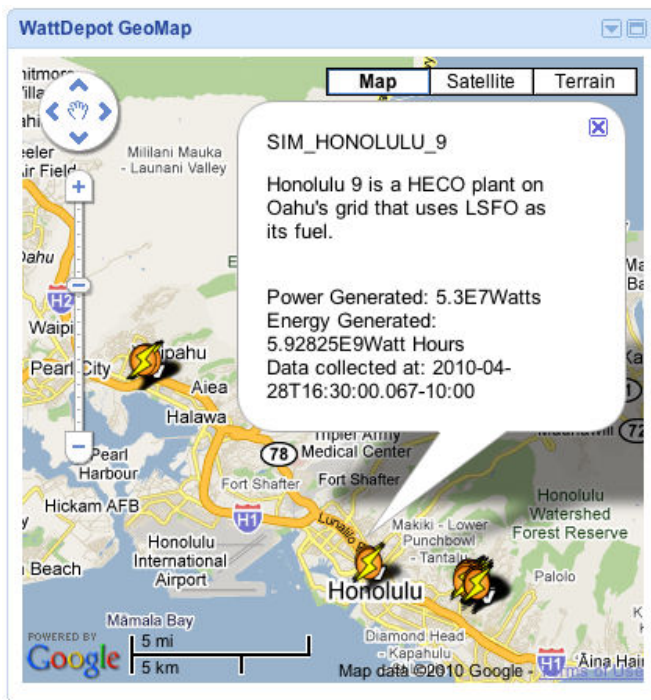


Fig. 4. *The WattDepot GeoMap client displaying generation sources for power production on Oahu, with simulated data for one plant displayed.*

one or more kinds of data to display, a time interval, a sampling interval, and a presentation mode (chart or table). The system retrieves the requested energy data from the WattDepot server. This application illustrates the utility of both virtual sources and automatic data interpolation. First,

the top blue trend line represents the power consumption of a virtual source aggregating together the power from the lower orange and yellow lines representing physical sources. Second, the sampling period of one hour does not correspond to the times at which energy data from the meters was provided. In fact, this visualization does not even require the user to know whether sources are virtual or physical, or when data is sent.

These three examples provide only a sampling of the kinds of WattDepot clients that are possible or currently under development. For example, we are developing an Android smartphone application that can display WattDepot data sources on any Android phone, and use it's geolocation facilities to let you know which energy Sources you are currently near.

## III. APPLICATIONS

Although WattDepot has been publicly available for only a short period of time, it is already finding use in a variety of application domains.

First, we are using WattDepot as infrastructure for an ambitious, next generation dorm energy competition web application. This application, called the Kukui Cup, combines traditional energy competition standings (based upon overall energy reduction) with both fine-grained, near-real time energy feedback and a parallel competition in which students carry out "energy literacy" activities to gain "Kukui Nut" points and compete for prizes based upon their point total. WattDepot will provide the underlying energy data collection, analysis, and visualization services for Kukui Cup.

Second, we have found WattDepot to be very useful as a simulation/prototyping mechanism. For example, we have been partnering with our local utility provider (Hawaiian Electric Company) to explore future customer-facing services. As part of this effort, we built a simple simulation of the

power grid on the island of Oahu in which we simulated meters collecting energy generation information from each of the 18 power plants on the island. By combining this information with data about the carbon intensity of each plant (provided by the Carma site [13]), we were able to develop a "stoplight" visualization similar that provided by the U.K. Ecotricity site [14]. Essentially, this visualization determines the carbon intensity associated with the current set of active power plants and their simulated current energy generation level on Oahu, and can inform consumers of whether this intensity is relatively low, average, or high.

Third, we are working with the Hawaii Natural Energy Institute (HNEI) on ways to use WattDepot as a data storage and normalization mechanism. HNEI is currently receiving regular data sets from affiliates with energy consumption data, and storage and manipulation of such datasets are already becoming unwieldy. Rather than build yet another one-off database solution, HNEI is evaluating WattDepot as a means to provide a master data repository and also to solve the "timestamp problem" by providing exported datasets with common timestamps across sources for input into their analysis tools.

## IV. Future Directions

In the near future, we plan to work on a number of promising research and development directions with WattDepot.

First, we want to explore the issues associated with privacy of energy data. Currently, WattDepot has a relatively simple privacy model. All sensors sending data for a particular Source must have the username and password corresponding to that Source's owner and their data requests are authenticated before being stored on the WattDepot server. This prevents unauthorized users from faking energy data. Each WattDepot Source must be declared as public or private. If public, then anyone can retrieve that data; if private, then only the Source's owner can retrieve the data.

We recognize that this privacy model is quite limited, and in fact all our applications to date have used public Sources, allowing their data to be freely accessed. We want to investigate more sophisticated and useful privacy models in future.

Second, we are interested in exploring alternative database back-ends, such as NoSQL systems like CouchDB, or cloud-based storage services like Amazon S3, Windows Azure, and so forth. The pluggable WattDepot architecture will enable us to experiment with different implementations and carry out performance analysis to assess their costs and benefits.

## V. Acknowledgments

## References

[1] "Google PowerMeter," http://www.google.com/powermeter.
[2] "Microsoft Hohm," http://www.microsoft-hohm/.
[3] "Smart energy 2.0," http://www.zigbee.org/SmartEnergy.
[4] "GridRouter Ecosystem," http://blog.smartsynch.com/tag/gridrouter-ecosystem/.
[5] "IEEE SCC21 Smart Grid Standards," http://grouper.ieee.org/groups/scc21/dr_shared/2030/.
[6] "Oasis Blue," http://www.oasis-open.org/resources/white-papers/blue/.
[7] "Control4 EMS 100," http://news.cnet.com/8301-11128_3-20000731-54.html.
[8] "The Energy Detective (TED) Home Energy Meter)," http://www.theenergydetective.com/.
[9] "Open Source Home Area Network (OSHAN)," http://sourceforge.net/projects/oshan/.
[10] "Open Phasor Data Concentrator (OpenPDC)," http://openpdc.codeplex.com/.
[11] "Representational State Transfer (REST)," http://en.wikipedia.org/wiki/Representational_State_Transfer.
[12] "WattDepot Application Programming Interface," http://code.google.com/p/wattdepot/wiki/RestApi.
[13] "Carma: Carbon Monitoring for Action," http://www.carma.org.
[14] "Ecotricity UK grid live website," http://www.ecotricity.co.uk/about/live-grid-carbon-intensity.