

LEARNING EMPIRICAL SOFTWARE ENGINEERING USING SOFTWARE  
INTENSIVE CARE UNIT

A THESIS SUBMITTED TO THE GRADUATE DIVISION OF THE  
UNIVERSITY OF HAWAI'I IN PARTIAL FULFILLMENT  
OF THE REQUIREMENTS FOR THE DEGREE OF

MASTER OF SCIENCE

IN

INFORMATION AND COMPUTER SCIENCES

DECEMBER 2009

By

Shaoxuan Zhang

Thesis Committee:

Philip M. Johnson, Chairperson

Henri Casanova

Scott Robertson

We certify that we have read this thesis and that, in our opinion, it is satisfactory in scope and quality as a thesis for the degree of Master of Science in Information and Computer Sciences.

THESIS COMMITTEE

---

Chairperson

©Copyright 2009

by

Shaoxuan Zhang

To myself,  
Shaoxuan Zhang,  
the only person worthy of my company.

# Acknowledgments

I want to “thank” my committee, without whose ridiculous demands, I would have graduated so, so, very much faster.

# Abstract

Abstract goes here.

# Table of Contents

Acknowledgments . . . . .	v
Abstract . . . . .	vi
List of Tables . . . . .	ix
List of Figures . . . . .	x
1 Introduction . . . . .	1
1.1 The Problem . . . . .	1
1.2 Software Intensive Care Unit Approach . . . . .	1
1.3 Thesis Claim . . . . .	1
1.4 Evaluation . . . . .	1
1.5 Thesis Structure . . . . .	2
2 Related Work . . . . .	3
2.1 TSP/PSP . . . . .	3
2.2 Research Based on Automated Data Collection . . . . .	4
2.2.1 Project ClockIt and Retina . . . . .	5
2.2.2 PROM . . . . .	6
2.3 Previous Case Studies of Hackystat . . . . .	7
3 Hackystat . . . . .	10
3.1 Hackystat Framework . . . . .	10
3.1.1 Sensors . . . . .	10
3.1.2 SensorBase . . . . .	11
3.2 Analysis Services . . . . .	12
3.2.1 Daily Project Data Analysis . . . . .	12
3.2.2 Telemetry Analysis . . . . .	12
3.3 Project Browser . . . . .	13
4 Design and Implementation of Software ICU . . . . .	14
4.1 Vital Signs . . . . .	14
4.2 Vital Sign Presentation . . . . .	18
4.2.1 Stream Trend Coloring . . . . .	19
4.2.2 Participation Coloring . . . . .	19
4.3 Mini Chart Drill-Down . . . . .	20
4.4 Vital Sign Configuration . . . . .	21
4.5 System Customization . . . . .	23
5 Classroom Evaluation . . . . .	25
5.1 Methodology . . . . .	25
5.2 Experimental Limitations . . . . .	25

5.3	Feedback regarding Hackystat system . . . . .	26
5.4	Verification of System Usage . . . . .	27
5.5	Utilities of Vital Signs . . . . .	29
5.6	Feasibility in a professional software development context . . . . .	31
6	Contribution and Future Directions . . . . .	32
A	2008 Classroom Evaluation Questionnaire of Hackystat . . . . .	33
B	Results form 2008 Classroom Evaluation Questionnaire of Hackystat . . . . .	36
	Bibliography . . . . .	53



# List of Tables

Table

Page

# List of Figures

<u>Figure</u>	<u>Page</u>
2.1 Progression of PSP . . . . .	4
2.2 ClockIt BlueJ Data Visualizer summary . . . . .	5
2.3 Data viewers of Retina. . . . .	6
2.4 Architecture of the PROM system . . . . .	7
2.5 Screenshot of course project to date analysis of Hackystat in 2003 . . . . .	8
2.6 Screenshot of file-metric telemetry analysis of Hackystat in 2006 . . . . .	9
3.1 Screenshot of file-metric telemetry analysis of Hackystat in 2006 . . . . .	11
4.1 A screenshot of Software ICU . . . . .	15
4.2 The Vital Sign Configuration panel in Software ICU . . . . .	22
5.1 The count of days when SICU was used along with the total invocation on per student basis. Each pair of c	
5.2 Analysis count on a per-student basis during the evaluation period. Each pair of columns represents data o	
5.3 Counts of selections of each vital sign in responses of question “If you used the Software ICU, please che	
5.4 Usage of Telemetry Analyses . . . . .	30

# Chapter 1

## Introduction

### 1.1 The Problem

Software metrics are the essential tool for performance measuring and quality control. But they are not comprehensive enough to make judgement with only one or two. So, we need to utilize multiple metrics in order to acquire insight of the health state of the software project. The more software metrics, the more comprehensive insight, but also the more effort is required to collect measure data and interpret values. My research tries to overcome this challenge by developing a system to help observe the health status of the software project from multiple software metrics in an effective way.

### 1.2 Software Intensive Care Unit Approach

What is SICU

### 1.3 Thesis Claim

We claim it is the most powerful tool ever in the world! =P

### 1.4 Evaluation

We evaluate in a classroom.

## **1.5 Thesis Structure**

Here is a paragraph that saying the same thing as TOC.

# Chapter 2

## Related Work

This chapter presents some works relate my research.

First part discusses previous research on empirical software engineering concepts. Most previous researches on measurement-based software engineering focus on methodology. Effective approaches are developed and deployed in actual practice. However, the lack of automation adds significant overhead to developers, thus lead to the impression that they are hard. Research of Hackystat and Software ICU is towards the 3rd generation of approaches to PSP metrics that automate data collection and analyze[9].

Second part discusses three recent research that focus on automated data collection. Two of them mainly focus on introductory level programming course and not very suitable to senior software development or professional settings. The third one is very similar to Hackystat and has related industry studies.

Last part discusses two previous related case studies of Hackystat system to provide some insight into the development history of Hackystat.

### 2.1 TSP/PSP

The Personal Software Process(PSP)[5] and the Team Software Process(TSP)[6] are among the most extensively studied approaches for measurement-based software engineering. They were developed by Watts Humphrey to teach students (in university and industry alike) the use of large scale methods based on the Capability Maturity Model (CMM). They scale down industrial software practices to fit the needs of small scale program development. Software processes and software engineering disciplines are gradually introduced through small program projects (e.g. course assignment projects). The PSP maturity progression is shown in [Figure 2.1](#). Students gather both process

and product measures of their projects. By comparing the measurement result to their original planning, they comprehend their programming habits, both pros and cons, and refine their process to higher level of maturity.

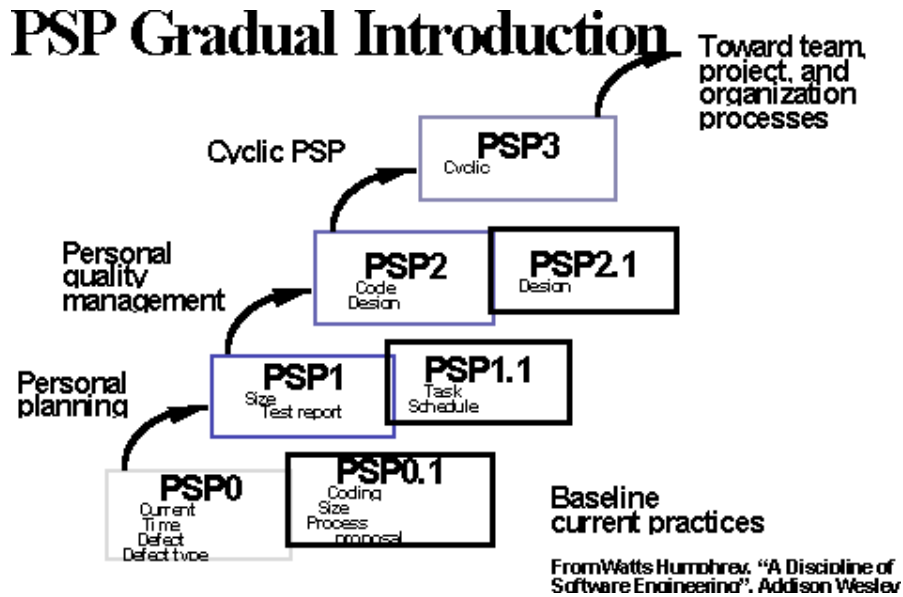


Figure 2.1. Progression of PSP

Their major drawbacks are lack of automation. Developers have to manually record their process and product data (mostly the development time and number of defects). The high overhead of data collection raises a barrier to its introduction and adoption. Additionally, it is not easy to “digest” the data. Developers have to manually analyze their logged data in order to understand their performance, then be able to improve it.

On the contrary, Software ICU provides a higher level of automation in tracking and analyzing software process and product data.

## 2.2 Research Based on Automated Data Collection

Project ClockIt and Retina are two recent research based on automated data collection to support entry-level programming courses, while PROM is the most similar research to Hackstat.

## 2.2.1 Project ClockIt and Retina

Project ClockIt provides a data logger as BlueJ<sup>1</sup> extension. It logs developer's open/close of project and package, file change and delete and compilation results. Data is logged to local file and later sent to a database via Internet. A data visualizer integrated into BlueJ is available to view data about the current project. Figure 2.2 shows an example of this visualizer. Data stored in database is used for statistic analysis such as class averages. A web interface is also available to instructors to view the individual data of their students and class average analysis data.





Figure 2.2. ClockIt BlueJ Data Visualizer summary

Closely related to ClockIt, Retina also provides automated data collection. Though Retina provide more tool support (BlueJ, Eclipse and command-line compiler), it focuses on a even smaller area of programming events: compilation. It gather data from students' compilation events, mostly compilation errors. In additional to its data viewer (see Figure 2.3), it also provides a recommendation tool for students. The tool use instant messaging (IM) to give students recommendation of approximated time needed for the upcoming assignment, and the compilation errors one is likely to make. These are based on both the student's previous data and the data from courses of previous semesters.

<sup>1</sup>"BlueJ is an integrated Java environment specifically designed for introductory teaching." –Quoted from <http://www.bluej.org/about/what.html>



(a) Retina Instructor Viewer

(b) Retina Student Viewer

Figure 2.3. Data viewers of Retina.

The difference between these two research and Software ICU is that they focus only on introductory level courses, where compilation is the interesting development event. In other hand, their relatively easy configuration overcomes one of the major short-coming of Hackstat and Software ICU. While neither of them provide good extensibility, they are unlikely to be useful in advanced programming situations like advanced programming course or professional setting.

## 2.2.2 PROM

PRO Metric (PROM) [11] is a system that similar to Hackstat. PROM is a software system for collecting process and product metrics in a software company. It is initiated and driven by the demand of the company, and thus the research is focus on industry setting. It is designed to work fully automatically without any interaction with the user in order to get reliable and accurate data about company internal workflows and development processes. It is organized in a sequence of interconnected components, communicating using SOAP protocol. Similar to sensors in Hackstat, it has plugins for many different applications, including IDEs, word processing tools, email clients, and issue tracking systems. Data then is transit to plugin server to extract metric, then the results are sent to PROM server to store into database. Figure 2.4 shows the overview of PROM's architecture.

PROM categorize users into 3 roles, developer, team leader, and manager of the team. Each of these role is provided different views of data. Developers have access to their individual, detailed data, the leader to the aggregated data of the whole team, and the manager to project level aggregated data.






Figure 2.4. Architecture of the PROM system

Compared to Hackstat, PROM’s data is stored as analyzed metrics result while Hackstat stores the raw sensor data. Different data viewers are provided to different groups of users while Hackstat provides same viewers to all users.

A case study of PROM in industrial environment[1] discusses the lessons learn from two year experience of using the PROM system in the IT department of a large company in Italy. Evidence indicates that adopting the PROM system requires a long set-up phrase and need the company and develop team’s patience and commitment to succeed, but it eventually delivers value to the company.

One of the lessons suggest that data presentation is as important as data accuracy and simplicity, brevity and clarity is preferable. Another lesson suggest that fast aggregated view of data is desired, and users of different roles favor different aggregations, e.g. developers like reports of their daily activities, while team leader and manager like summary views of data on team and project level. Software ICU’s simple and fast data presentation and high configurability and extensibility is suitable to address these requirements.

## 2.3 Previous Case Studies of Hackstat

The classroom study presented in this thesis is the third case study of Hackstat system in a classroom setting.

The first case study performed in 2003 used an early version of Hackstat[7]. During that time, Hackstat was only collecting 4 types of metrics (Active Time, Size, Unit Tests and Coverage). The system was oriented around a set of “Course” analyses that were tailored to an educational setting. Those analyses summarize the individual team project metric data in tabular form, and also presented comparisons of all of the course projectsFigure 2.5. The evaluation show

that the installation of Ant sensors is the most significant barrier to the system. It was too difficult to install without direct help from the development team. But the overhead of use is relatively low and analyses were usable and useful. Data privacy was uncomfortable for some students.

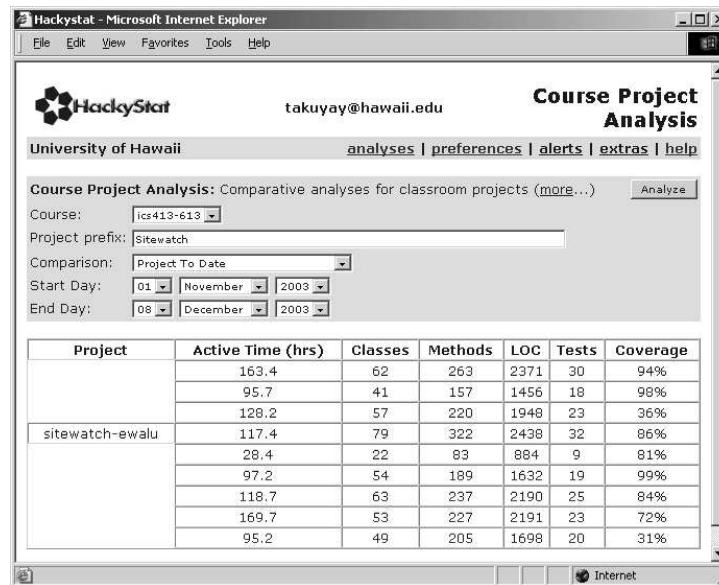


Figure 2.5. Screenshot of course project to date analysis of Hackystat in 2003

The second case study in 2006 is a partial replication of the first case study[8]. Hackystat had undergone significant change from 2003 to 2006. The sensor installation, which is the major barrier to the system in 2003, was automated by the hackyInstaller GUI, which greatly lower the overhead of configuration for developers. The evaluation also shows significant drop in sensor installation difficulty. However, new sophisticated Telemetry analysis [Figure 2.6](#) and its complex user interface raise the difficulty of using it and interpreting data, lead to slight drop in usability and professional feasibility.

In 2007, Hackystat was re-implemented using a new architecture. Adopting service-oriented architecture enable multiple user interface separate from the services components. The Software ICU is built upon a new web-based UI called Project Browser, and the classroom study is also based on this user interface.




Figure 2.6. Screenshot of file-metric telemetry analysis of Hackstat in 2006

# Chapter 3

## Hackystat

In this research, I utilize Hackystat to implement the Software ICU. This chapter briefly introduce the Hackystat system, which was invented by Professor Philip M. Johnson, in the Collaborative Software Development Laboratory, Department of Information and Computer Sciences, University of Hawaii at Manoa.

### 3.1 Hackystat Framework

Hackystat is an open source framework for collection, analysis, visualization, interpretation, annotation, and dissemination of software development process and product data. Hackystat consist of many software services that communicate using REST architectural principles[?]. These software services can be categorize into 4 groups, sensors, data repository, analysis services and viewers. [Figure 3.1](#) shows the architecture of Hackystat system.

#### 3.1.1 Sensors

Sensors are small software plugins that collect data from the use of tools and applications. Currently, sensors are available in many development software including Eclipse, Emacs, Ant, etc. Sensor data is represented in XML, and consist of seven basic elements: data owner, resource, timestamp, runtime, tool, Sensor Data Type and properties. The first six are required and the last one is optional.

Sensor Data Type(SDT) is specified in every piece of sensor data when collected, so that the same type of data can be collected from different tools and higher level services can easily determine which data is relevant to them. Sensor data is designed to record only an piece of atomic data such as size of a single file, and runtime is use to group data that belongs to the same event,

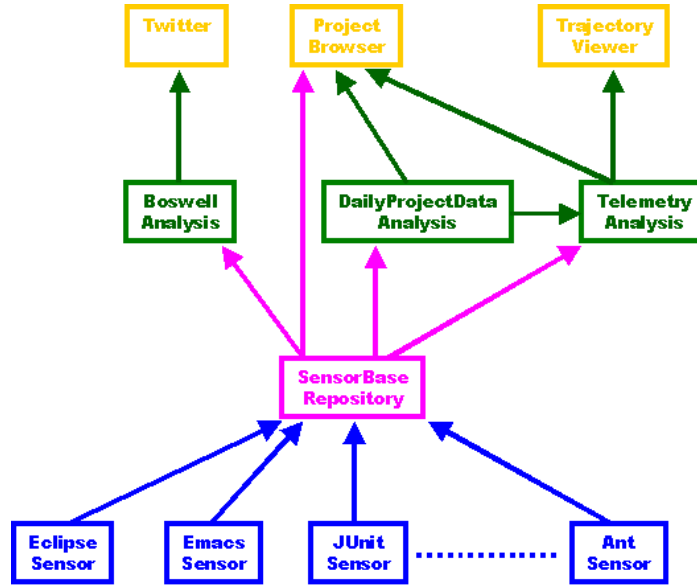


Figure 3.1. Screenshot of file-metric telemetry analysis of Hackstat in 2006

such as the file metric of a project. Properties are additional information for different types of sensor data, such as coverage value for coverage SDT and lines of code for file metric.

Sensors are designed to work automatically without any attention of user besides initial configuration. In order to reduce internet communication and support offline work, data is temporarily stored, then sent to data repository every several minutes or when internet connection is available.

### 3.1.2 SensorBase

Sensor data is sent to the data repository, called SensorBase. SensorBase store the data as it sent from sensors, and provide RESTful interface for easy use. Sensor data can be queried with the six required elements mentioned above via HTTP calls, and data is sent back as XML. SensorBase is implemented with a database manager abstract class, thus it is easy to add support to different database implementations. Current version of Hackstat provides database support for Derby, Oracle and PostgreSQL.

## 3.2 Analysis Services

Analysis services of Hackystat provide abstractions of the raw data from SensorBase. DailyProjectData and Telemetry are the two most fundamental analyzes of Hackystat.

### 3.2.1 Daily Project Data Analysis

As its name says, DailyProjectData(DPD) service provides abstractions of sensor data associated with a single project within a 24 hour period, which represents a simple software development metric of a single day. Data of a single project includes data from all members of that project. In a DailyProjectData instance, both summary value, e.g. total development time across the project, and detail values, e.g. total development time of a project member, are available. So it is easy for higher level service to use this data.

Each DPD analysis generates software metric from data of a certain Sensor Data Type. Current available DPD analyses are Build, Code Issue, Commit, Complexity, Coupling, Coverage, Dev Time, File Metric, Issue, and Unit Test. These DPD analyses are the basic of the Hackystat system, all other analysis services are based on DPD. While DPD is the lowest level of abstraction, these can also be considered as the available software metrics in Hackystat.

### 3.2.2 Telemetry Analysis

Based on DPD service, Telemetry service provides abstraction over a longer period of time such as several days, weeks or months. A Telemetry Chart consists of one or more streams of data points. Each data point represents the metric value of in a single granularity (day, week or month). Together they show the trend of the metric(s).

There is a special group of Telemetry charts called Member-Level Telemetry. These charts consist of several stream, each of which belongs to a project member. They are used in Software ICU's drill down feature to compare performance of each member within a project.

To support the work practices of different organizations, Telemetry service provides a domain specific language that allows to build new Telemetry Chart with Telemetry stream lines. The predefined Telemetry Charts are all written with this language.

Telemetry streams can also accept parameters to refine the object data. This feature is inherited in Software ICU, where user can configure the parameters of each Telemetry analysis of each vital sign (more detail discuss in [section 4.1](#) and [section 4.4](#)).

### 3.3 Project Browser

Project Browser is one of the viewers in Hackystat system. It is based on Wicket<sup>1</sup>, a Java-based web application framework. Project Browser is integrated with viewers to all Hackystat services, which are organized as tabs.

With help of Wicket's modularization, viewers on Project Browser can share many common panel, such as project/date selection panel and Ajax loading process panel, which facilitate the development of new page. This also makes user's experience more consistent across different viewers. Therefore it now serves as a data presentation and high level analysis development center. Several new presentations and high level analysis are developed upon it, Software ICU is one them.

---

<sup>1</sup><http://wicket.apache.org/>

## Chapter 4

# Design and Implementation of Software ICU

In order to achieve good governance of software development projects, we adopt the metaphor from medical ICU and develop a system called Software Intensive Care Unit. It consists of a set of vital signs, each of which is based on one software development metric and indicates the project's "health" state from one perspective.

Interface of Software ICU is separated into two part. The left-hand side is the control panel, where user can pick the analysis period, data granularity and select projects to analyze. The right-hand side consists of three panels: the data panel, the loading process panel, and the configuration panel. Each of these panels is discussed in following sections. Data panel is the major panel to show the result of SICU analysis. [Figure 4.1](#) shows an example of Software ICU.

### 4.1 Vital Signs

Similar to medical ICU, using multiple software development metrics in Software ICU are necessary because there is not a single metric can determine the health state of a software project. Similar to medical vital signs, each software metric show the process or product state from different aspect. And changes in one of them may or may not indicate a change in health state, but changes in more of them indicates higher possibility that health state changed. In this study, we use nine vital signs in Software ICU.

Vital signs of software projects are measured by various software development process or product metrics. Each of these vital sign reveal an aspect of the health state of the software project. In this section I will discuss all these vital signs.





Figure 4.1. A screenshot of Software ICU

**Coverage** Coverage is a good indicator of the tests' quality. It stands for the test coverage of source code in unit testing, which usually measured as the percentage of code units (line, method, class, etc) that is executed during unit test. There are a number of coverage criteria, such as line, method, class, conditional, etc. In Software ICU, user can select which to use. No matter which criteria to use, higher coverage is always better because higher percentage of code covered by unit testing indicates lower risk of existence of bug in untested code segment. However, high coverage does not necessary mean good quality of unit test, and vice versa. One of the reason is that, in some situations, it is difficult to achieve high test coverage because of the difficulty of verifying results, especially when using UI frameworks. Another reason is that the code executed during unit testing can be unverified. For example, when testing an image processor with a given image file, the code of loading the image file is executed, but the test does probably not have assertion about the correctness of loading the file. But as long as developers don't have the intend to trick coverage in order to pretend to be writing enough unit tests (It is possible if coverage is used to judge their performance.), raising coverage is always a good thing.

**Cyclomatic complexity** Cyclomatic complexity, a measurement of the complexity of a program developed by Thomas J. McCabe, measures the number of linearly independent paths through a program's source code[10]. The higher cyclomatic complexity, the more distinct control paths in a program module, and the more difficult to achieve high path test coverage. Additionally, code of high complexity is difficult to understand, thus it is hard to maintain. Therefore, program modules are preferred to have lower complexity. But high complexity is not necessary evil. The nature of some programs just require high level of complexity. Also raising in complexity is sometimes unavoidable during development, especially when optimizing code performance. However, developers should try to lower complexity, especially in early stage of development, so that it is easier for future maintenance.

**Coupling** Coupling, or dependency, is the degree to which each program module relies on each one of the other modules[2]. It is a measurement of the complexity of the whole system's module reference tree. Whenever one module is modified, there will be a chance that the changes may cause bugs in one of modules that relies on it. Therefore, higher coupling implies higher risk of introducing bugs when making changes, thus harder to maintain. High coupling might also be harder to reuse because dependent modules must be included. Therefore, Coupling is suggested to be kept low.

**DevTime** DevTime, abbreviation of Develop Time, is a measurement of develop time of developers. Hackystat use a special approach to measure this: for each 5 minutes interval, if there is anything development activities happen, the developer is consider developing in that interval. It relax the criteria of measuring develop time so that coding while reading from documentation will get the same DevTime as intensive coding period. Thus it is more accurate than previous ActiveTime in older Hackystat. However, Hackystat sensors for DevTime are only available in several IDEs (currently available to Emacs, Eclipse, and Visual Studio). No sensors are available for other applications that might be used during developing, such as browsers, E-mail clients, office systems, or other editors/readers. So the monitored development activities is limited. Moreover, some developing activities, such as reading and learning, are very difficult to track. Therefore, DevTime should not be simply used to determine a developer's effort. But as habit of an individual will not change a lot, so the DevTime of a developer should be relatively stable over time. Thus large sudden increase in DevTime is a possible sign of bad developing habit like "start late near deadline".

**Churn** Churn is a measurement of the change (add, delete and/or modified) of code that is made into repository. It is usually measured by LOC (lines of code). It is an indicator of developers' contribution to the project. Interpretation of this metric is conditional. In the early stage of development, churn is expected to be higher because new codes are keep being added. During the maintenance of a system, churn is mainly from fixing bugs and adding new features, both of which are fewer for a stable system, thus churn is expected to be lower. In terms of develop behavior, churn of a developer reflects the amount of work of him/her. It tends to be relatively stable over time in the same project because the work rate of an individual does not vary a lot in the same coding condition. Dramatically change in churn of an individual developer while DevTime not changing respectively is a bad phenomena, which might due to bad developing habit like "copy and paste without understanding".

**Commit** Commit measure the number of commitments made into repository. "Commit early, commit often" is one of the well-accepted guideline of continuous integration. For the same amount of churn, more commits implies better following of this discipline.

**Size** The size of the project is measured by the source lines of code (SLOC), which counts the number of lines in the text of the program's source code. It can be a sign of the effort put into the project, However, SLOC alone does not make as much sense about the state of the project as Churn. We include this vital sign only to give user an idea of the size of the project, just like the height in your medical record.

**Test** Test is a count of unit test tasks invoked in a period of time. Unit testing is a software verification and validation method in which a developer tests an individual units of source code. It is used to ensure that code meets its design and behaves as intended. A requirement of good development behavior is to test while coding, or event better, use "Test Driven Development" (TDD). No matter what development pattern to follow, unit testing is a dispensable component and regular execution of unit tests is always a good sign of "health" development habit.

**Build** Build is a count of ant build task invoked in a period of time. A build task accomplish all necessary step to ensure the correctness of the code before commit. It consists of compilation, code inspection, unit testing, documentation generation, etc. It is an usual activity in software development nowadays. Though how often to build largely depends on personal preference and habit, it is advised to build often to ensure the correctness of the system.

These nine vital signs are the default set in Software ICU, but not the only setting. Users can determine to use which vital signs in their setting, as well as creating new vital sign analysis with Telemetry charts. More detail about these configuration and customization are discuss in [section 4.4](#) and [section 4.5](#).

## 4.2 Vital Sign Presentation

As reported in case study of PROM, data presentation is as important as data accuracy[1]. One of our primary goal of Software ICU is to provide a proper presentation to help interpret large amount of software metrics data. In order to achieve this goal, Software ICU use mini charts to integrate historical data and use color to categorize the health state of a vital sign.

A vital sign analysis consists of two part: a numerical latest value and a mini historical chart.

**Latest Value** represents the newest state of the vital sign in the analysis period. In our implementation, it shows the most recent associated DailyProjectData. If there is no DPD on the latest date of the analysis period, it will search back the time period for the first available data of that DPD. The latest value will be “N/A” only when there is no any data of that metric in the whole analysis period.

**Mini chart** represents the trend line of the associated metric data over the analysis period. This mini chart is implemented as bar charts. Each bar represents the DailyProjectData value of the metric on a unit of granularity (day, week or month). Bars heights are scaled so that the highest bar is almost reach the top of the chart.

However, only use the last values and mini charts does not address the requirement of fast data interpretation. Thus we further enhance the representation by adding colors to those numerical values and charts to provide intuitive idea of the “health” state of the vital signs.

Generally, we use color green to represent “healthy” state, red to “unhealthy” state, and “yellow” to state between green and red or uncertain. This is color pattern is good for indicating states because it match the convention people comprehend color and thus most people can understand it without reading instructions.

Different vital sign may use different coloring method, and the latest value and the mini bar are colored separately. Choice of coloring method mainly depends on the nature of the vital sign. In general, vital signs that have clear preference of higher or lower, like most based on software

development product metrics (Coverage, Complexity, Coupling) will use Stream Trend Coloring method, and vital signs based on software process metrics will likely to use Participation Coloring method. Sometimes, there may be no ideal coloring method for a vital sign, such as FileMetric, then user can select to leave that vital sign uncolored.

#### 4.2.1 Stream Trend Coloring

Stream Trend Coloring method determines the health of a metric by its value and trend. It colors the latest value as well as the mini chart. It takes three parameters: *HigherBetter*, *Higher Threshold* and *Lower Threshold*. User can decide the preferable trend, higher or lower, using the *HigherBetter* parameter. A raising mini chart is considered to be good/bad if the *HigherBetter* parameter is set to true/false. A trend is considered to be raising if there is no any value point is lower than the one before, and the last value is greater than the first value. Falling trend is considered the same opposite way. In order to be able to categorize trends that have some small disruption as raising or falling, the Stream Trend Coloring method consider small amount (proportionate to the average of the first and the last value) of change as equal. Stable trends are always considered as “healthy” because in that case it is as good as “healthy” that user doesn’t need to pay much attention to it, while the actual value will be shown in the latest value where the value will be judged to be “healthy” or not. And unstable trend is marked as yellow because it is no fast way to tell if it indicates a good state or not.

*Higher Threshold* and *Lower Threshold* parameters are only used when coloring the latest value. Values exceeds the higher threshold will be colored green if *HigherBetter* is true, or red if *HigherBetter* is false. Values lower than the lower threshold are colored in similar way. Values between these two thresholds are always colored yellow.

#### 4.2.2 Participation Coloring

Participation Coloring method determines the health of a stream by the participations of all members in the project. It only colors the mini chart, leaving the latest value uncolored. This coloring method is designed to detect the health state of team cooperation, mainly via software process metrics. It takes three parameters: *Member Percentage*, *Threshold* and *Frequency*. Participation Coloring method colors a mini chart green if

1. there are more percentage of members than the *Member Percentage* parameter that,
2. have the metric value greater than or equal to the *Threshold* parameter per day,

3. for more frequently than the *Frequency* parameter in the analysis period.

A mini chart is colored yellow if it does not meet the green requirement, but the metric of the team as a whole meet the requirement of green, i.e.,

1. the combined metric value is greater than or equal to the *Threshold* parameter per day,
2. for more frequently than the *Frequency* parameter in the analysis period.

If the yellow requirement is not meet neither, the mini chart will be colored red.

In other words, Participation Coloring method color a vital sign green if most members of a project are making noticeable contribution to the project regularly, and color it yellow if not but there is someone making contribution to the project in most of the time, and color it red otherwise, which means, in terms of this vital sign metric, the contribution of members to the project is rare and/or insignificant.

### 4.3 Mini Chart Drill-Down

In each non-empty mini chart, Software ICU provides a link of the drill-down Telemetry analysis. The drill-down Telemetry analysis is the analysis that used to generate the mini chart. For most of software product metrics, such as Coverage, Complexity and Coupling, the drill-down Telemetry analysis will show the same chart as the mini chart in Software ICU's vital sign block, just in different style with more detailed axes. However, for software process metrics, instead of the original chart, an associated member-level Telemetry analysis is shown when drill-down.

The member-level Telemetry shows multiple stream lines in the chart, each of which represents the analysis of a member of the project. From this member-level Telemetry, it is easy to see members' participation to the project in this metric. This is most useful when combined with the Participation Coloring in Software ICU, where you see the summary result of members' participation, and then understand the detail with member-level Telemetry analysis.

The drill-down Telemetry analysis use the same parameters as used in Software ICU, thus the non-member-level chart should be identical with the one in Software ICU. Vital signs with drill-down to member-level Telemetry, use as well the member-level Telemetry to generate the mini chart in vital sign block by summing all stream lines into one. Software ICU provides different integrating method to handle vital signs that use member-level Telemetry, more detail is discussed in [section 4.5](#).

## 4.4 Vital Sign Configuration

Software ICU provides user full access to the configuration of each vital sign. User can enable/disable a vital sign, choose its coloring method, and configure the parameters of the associated Telemetry analysis.

By clicking the “Show Configuration” button in the input panel on the left, user can open the configuration panel to its right. Then the “Show Configuration” will be disabled when the configuration panel is shown. **Figure 4.2** shows an example of the vital sign configurations. The first column is the name of the vital sign with the checkbox to enable or disable a vital sign. When a vital sign is disabled, the configuration of color method and Telemetry parameters will be disappear, however, the settings are not discarded, thus when enabled again, it be the same before disabled. The second column is the color method. Current version of Software ICU provide three choices: StreamTrend, Participation and None. By choosing the first two, its associated parameters, which discussed in **section 4.2**, are shown next to the drop-down selection field. When “None” is selected, nothing will be shown in that space. The last column is the Telemetry parameters, which is defined in the definition of Telemetry charts, and will be directly transferred to Telemetry service when retrieving Telemetry analysis for vital sign presentation. Because of the implementation reason, results of enabling/disabling a vital sign and selecting different color method will be saved immediately, but other fields will only be saved when the “OK” button in the bottom is pushed. When the “OK” button is pushed, the configuration panel will disappear after setting is saved, and the “Show Configuration” button will become available again.

In order to persist user’s configuration setting between each visit, the configuration settings are saved in server side using UriCache. UriCache is a wrapper around the Apache JCS system<sup>1</sup>. It is designed to provide an API well suited to the needs of Hackystat services. The vital sign configuration objects are directly cached, under the name of the user. The cache expiration timer is set to 300 days so that it will not easily be expired. But if the cache is expired, the system will use the default setting of the vital signs.

Next to the “OK” button is the “Rest to Default” button. It will restore all vital sign configuration settings to default, and the result of restoring will be shown and saved immediately.

In the bottom-right of the configuration panel is a link called “Configuration Instructions”. When clicked, it will show a simple instruction of the configuration panel in a pop-up window.

---

<sup>1</sup>“JCS is a distributed caching system written in java. It is intended to speed up applications by providing a means to manage cached data of various dynamic natures.” –<http://jakarta.apache.org/jcs/>



## Software Project Portfolio Analysis

**From Date:** 2009-10-02 17  
**To Date:** 2009-10-06 17  
**Granularity:** Day  
[Show Configuration](#)

**Project(s):**  
☐ AmbientHackstat  
☒ Default  
☐ Hackstat  
☐ hackstat-analysis-d  
☐ hackstat-analysis-te  
☐ hackstat-sensor-ant  
☐ hackstat-sensor-ecl  
☐ hackstat-sensor-em  
☐ hackstat-sensor-exa  
☐ hackstat-sensor-ma  
☐ hackstat-sensor-sho  
☐ hackstat-sensor-vim  
☐ hackvstat-sensor-xm  
[OK](#) [Cancel](#)

Measure Name	Color Method	Telemetry Parameters
<input checked="" type="checkbox"/> Coverage	StreamTrend	Higher Threshold: 90 Lower Threshold: 40 Higher Better: <input checked="" type="checkbox"/> mode: Percentage granularity: method
<input checked="" type="checkbox"/> Complexity	StreamTrend	Higher Threshold: 0 Lower Threshold: 0 Higher Better: <input checked="" type="checkbox"/> mode: AverageCc threshold: 10 tool: JavaNCSS
<input checked="" type="checkbox"/> Coupling	StreamTrend	Higher Threshold: 20 Lower Threshold: 10 Higher Better: <input type="checkbox"/> coupling: All mode: Average type: class threshold: 10 tool: DependencyFinde
<input checked="" type="checkbox"/> Churn	StreamTrend	Higher Threshold: 900 Lower Threshold: 400 Higher Better: <input type="checkbox"/> cumulative: <input type="checkbox"/>
<input checked="" type="checkbox"/> Size(LOC)	None	sizemetric: TotalLines tool: *
<input checked="" type="checkbox"/> DevTime	Participation	Member(%) 50 Threshold 0.5 Frequency(%) 50 cumulative: <input type="checkbox"/>
<input checked="" type="checkbox"/> Commit	Participation	Member(%) 50 Threshold 1 Frequency(%) 50 cumulative: <input type="checkbox"/>
<input checked="" type="checkbox"/> Build	Participation	Member(%) 50 Threshold 3 Frequency(%) 50 result: * type: * cumulative: <input type="checkbox"/>
<input checked="" type="checkbox"/> Test	Participation	Member(%) 50 Threshold 10 Frequency(%) 50 mode: TotalCoun cumulative: <input type="checkbox"/>
<input type="checkbox"/> CodeIssue		

[OK](#) [Reset to Default](#) [Configuration Instructions](#)

Figure 4.2. The Vital Sign Configuration panel in Software ICU



## 4.5 System Customization

Beside the ability to configure vital signs on the fly, Software ICU provide also offline customization of default vital signs. All vital signs, including the default set discussed above, are defined in PortfolioDefinition XML files. There are two place the system will look for these XML files. The first place is inside the package of detail panel of Software ICU, where the default set of vital signs are defined. The other place is `/.hackystat/projectbrowser/`, where “ ” stands for user’s home directory. Here is an example of the definition XML file:

```
<?xml version="1.0" encoding="utf-8"?>
<PortfolioDefinitions>
  <Measures>
    <Measure name="Coverage"
      classifierMethod="StreamTrend"
      enabled="true"
      telemetryParameters="Percentage,method">
      <StreamTrendParameters higherBetter="true"
        lowerThresold="40"
        higherThresold="90"/>
    </Measure>
    <Measure name="MemberDevTime"
      alias="DevTime"
      merge="sum"
      classifierMethod="Participation"
      enabled="true">
      <ParticipationParameters memberPercentage="50"
        thresoldValue="0.5"
        frequencyPercentage="50"/>
    </Measure>
    <Measure name="FileMetric"
      alias="Size(LOC)"
      enabled="true">
    </Measure>
  </Measures>
</PortfolioDefinitions>
```

There is a root element called *PortoflioDefinitions*, enclosing a single element *Measures*. Within the *Measures* element, there are a set of *Measure* element, each of which stands for a vital sign. Each *Measure* element can take up to six attributes:

1. The *name* attribute is required. It is the name of the Telemetry analysis used in this vital sign.

2. The *alias* attribute is optional. When it is set, it will be used as the name of this vital sign. Otherwise, the *name* attribute will be used as this vital sign's name.
3. The *classifierMethod* attribute defines the default coloring method, either *StreamTrend* or *Participation*. This attribute is optional. When it is unset, the default coloring method will be none.
4. The *enabled* attribute defines if the vital sign is enabled by default. If set to false, the vital sign will be disabled by default. But user is still able to enable it in configuration panel.
5. The *merge* attribute defines the method to integrate multi-stream telemetry. It is necessary for member-level telemetries to work. "sum", "min" and "max" are available choices. If it is unset, the first stream of the telemetry will be used. Because of the order of streams in a multi-stream telemetry is not guaranteed, use member-level telemetry without setting this attribute might cause unexpected result.
6. The *telemetryParameters* attribute is the Telemetry parameters of the telemetry analysis defined in *name* attribute. It can be unset, then the default parameters will be used. This attribute accept value formatted the same way as Telemetry Rest API, i.e. common separated values ordered the same as parameter definition in the Telemetry analysis.

The *Measure* element also have up to two optional sub-elements. They are *StreamTrendParameters* and *ParticipationParameters*, each of which defines default parameters of the corresponded coloring method, and can exist together regardless what is set in the *classifierMethod* attribute. They take attributes same as their parameters discussed in [section 4.2](#).

# Chapter 5

## Classroom Evaluation

### 5.1 Methodology

The system is gradually introduced to the students in Software Engineering (ICS413) course on Fall 2008. At the end of the Fall 2008 semester, the students in ICS 413 were contacted by email and asked to respond to the a questionnaire soliciting their opinions regarding Hackystat and Software ICU. The complete questionnaire can be found in [Appendix A](#). I provided each of the students a “secret” code. The correspondences between the secret codes and the students are only known by me, but not the instructor of the class, in order to avoid the potential for bias to “please” the instructor/designer who would presumably be gratified by positive responses to the questionnaire. Response was optional, but the students were offered extra credit points for providing their opinions. The list of names who should be awarded extra credit was sent to the class instructor without identifying individual responses. The students were asked to reply within a week. Eighteen out of the nineteen students contacted provided responses.

In addition, students’ usage is logged in the system, which is not aware by students, and is compared to the feedbacks from the survey.

The results of the classroom evaluation questionnaire can be found in [Appendix B](#)

### 5.2 Experimental Limitations

Before drawing any conclusions from this data, it is important to recognize the limitations of this study. Compared to the limitations associated with previous study in 2003 and 2006, anonymity is achieved, but others are still unsolved in class evaluation.

First, this data is drawn from a limited sample size of 18 students in software engineering classes at the University of Hawaii. The subjects therefore have a relatively narrow and homogeneous background in software development.

Second, the context in which they used the system was a course project. Course projects tend to be smaller, narrower in scope, and with less pressure on the developers than an industrial context. It is one thing to get a poor grade for doing a poor job, it is another thing to lose your job for doing a poor job. In addition, students are not working full-time on the system; the development project is just one assignment among several.

These are all major limitations on the external validity of the responses. They do not make the results meaningless, but rather help provide a perspective on how to gain additional evidence in future that would confirm/disconfirm these initial findings. For example, it would be helpful to deploy Hackystat in a real software company, and then gather data anonymously from the coders and managers. Other insights into future research directions will be covered in an upcoming section.

### 5.3 Feedback regarding Hackystat system

Besides the purpose of research regarding Software ICU, this study is also regarded as a evaluation of Hackystat of the new architetur.

The responses of the questionnaire indicate that sensors installation is more difficult than the Hackystat in 2006. This is not surprised because of the fact that a client-side installer package is provided in 2006, which is not yet available in the time of this study. However, once the sensors are installed correctly, no further effort is required in data collection. Because all the students are using the public services of Hackystat<sup>1</sup>, there is no effort required in the server-side configuration, which is reported to be the biggest installation/configuration difficulty in Hackystat case study in 2006.

The sensors' installation difficulties is mainly cause by the documentation. Though installation guides are provided in every components, the documentation is too distributed to follow as a result of Hackystat's service-oriented architecture, which reduced the coupling among components, but also reduced the correspondence among components' documentation.

Regarding development data sharing, most students felt OK with sharing development data with other members. But three students had concerns that sharing development data would

---

<sup>1</sup>SensorBase: <http://dasha.ics.hawaii.edu:9876/sensorbase>, DailyPro-  
jectData: <http://dasha.ics.hawaii.edu:9877/dailyprojectdata>, Teleme-  
try: <http://dasha.ics.hawaii.edu:9878/telemetry>, ProjectBrowser:  
<http://dasha.ics.hawaii.edu:9879/projectbrowser>

reveal their programming habits and introduce too much competition of statical stats, which made them nervous. It is interesting that those three students are the three with least Software ICU running count in Figure 5.2. It is reasonable to infer that they worked less harder than other students and did not want to be noticed. However, this is an example of how measurement dysfunction might cause negative effect on developers.

## 5.4 Verification of System Usage

??, Figure 5.4 and Figure 5.2 show the data from system usage logging. I combine it with the data from questionnaire to confirm that the students' responses from questionnaire are reflecting the truths of their practice.




Figure 5.1. The count of days when SICU was used along with the total invocation on per student bias. Each pair of columns represents data of one student. The X axis shows the responses from questionnaire. VH = every day or more; H = 2-3 times per week; M = once a week; N/A = not available.

When verifying the questionnaire responses against the log data, we find that the choices of question “How frequently did you use the telemetry page?” and “How frequently did you use the Software ICU?” are somehow ambiguous. Though “every day or more” is surely asking how many days you use the analysis, “2-3 times a week” may be understood as times of invocations. Figure 5.1 shows data of these two interpretations. If we consider the answers as “days of use”, the actual use frequencies are much lower than reported, because there are 28 days in the evaluation period but the

highest number of days of use is only 18. But if we consider the answers as “times of invocations”, the invocation frequencies are more matched to reported frequencies. However, in either case, the difference of actual usage between students who claim to use SICU analysis “every day or more” and “2-3 times a week” is not obvious. Though the total invocation times and days of the first group is higher than the second, some students of the second group used SICU analysis more frequently than the students of the first group. But this error is acceptable because the frequency of use is just as remembered and might not be precise. So if the criteria is weakened and both “every day or more” and “2-3 times a week” are considered as “did use SICU frequently”, all responses match their log data except three of them. Those three students claimed that they use Software ICU 2 to 3 times a week or more, but actually only half as much as they claimed (The lowest one with response “every day or more” and the lowest two with response of “2-3 times per week”).

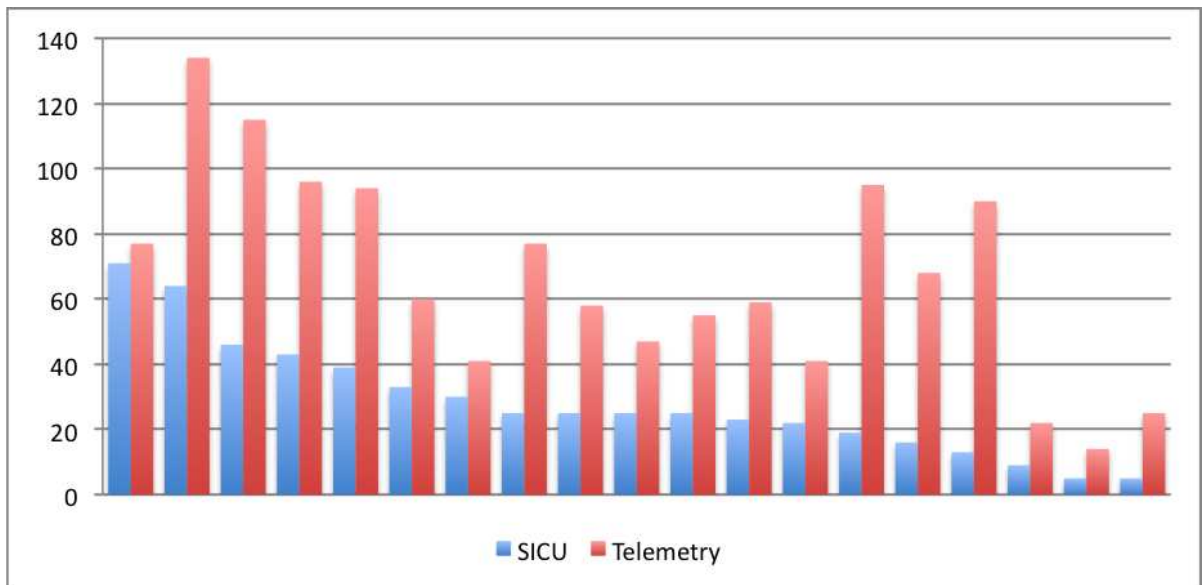


Figure 5.2. Analysis count on a per-student basis during the evaluation period. Each pair of columns represents data of one student.

I also find that though the reported frequency of SICU and Telemetry are similar, Telemetry’s analysis invocations are much more than SICU’s (see Figure 5.2). But this matches the native of these two analyses: SICU shows the overall summary of a project’s health and no need to run more than once a day, while Telemetry shows detail of a vital sign and would often be run multiple times in every use.

Because both questionnaire responses and log data show matched evidence that students are using Software ICU and other Hackstat services frequently, I believe that students participated

in this evaluations actually have plenty of use experience of Software ICU and Hackystat and the responses of the questionnaire are based on their real experience and opinions.

## 5.5 Utilities of Vital Signs




Figure 5.3. Counts of selections of each vital sign in responses of question “If you used the Software ICU, please check the vital signs that were useful to you.”.

Regarding Software ICU as a whole, 7 out of 9 vital signs are considered to be useful by at least half of the respondents (Figure 5.3). 10 out of 18 responses said Software ICU was accurately reflecting the health of their project via colors. Other 6 responses are not denying the utilities of vital signs, but are arguing that some vital signs are not accurate enough to determine a project’s health. Only one student found it is “hard to determine what will fall into green, red, or yellow”, and the last student said he failed successfully configure the sensors. Overall, students were quite positive regarding the utilities of vital signs.

Three vital signs are mostly concerned in students feedback: Coupling, Churn and Dev-Time.

## 5.6 Vital Sign Popularity




Figure 5.4. Usage of Telemetry Analyses

The number of invocations of Telemetry analyses can be an indicator of vital sign popularity and usefulness. Both log data(5.4(b)) and questionnaire responses indicate that Telemetry is mainly used to run member-level analyses. The two most used analyses are MemberDevTime and MemberCommit(5.4(a)). In the responses of the question of vital sign usefulness(Figure 5.3), DevTime and Commit are also among the three most popular vital signs. The other one in top-three is Coverage.

It is not surprised that Coverage is among the most popular vital signs. Compared to other productive metrics, it is the most intuitive indicator of project's quality. It is not in the frequently used Telemetry analyses is possibly because there is no need to run a separate Telemetry analysis. Users can get all the information from the coverage vital sign in Software ICU.

But DevTime and Commit's popularities are not expected prior to the evaluation. Survey result indicates that, this is not a special case: vital signs based on software process metrics attract much more attention than those based on software product metrics (Figure 5.3). Popularity of pro-



cess vital signs (DevTime, Commit, Build, Test, and Churn) exceed all productive vital signs except Coverage (Complexity, Coupling and Size). There are three major reasons that lead to this result.

The first reason is those popular vital signs are much easier to interpret than those are not popular. Meaning of these vital signs are very straight-forward. On the contrary, as mentioned in students comments, complexity and coupling metrics are more difficult to comprehend. Though the general guidelines of them are lower the better, the meaning of a certain number is not easy to understand because of the nature of these metrics. Also, as the the development progress and more features and functions are added to the code, complexity and coupling are always tend to increase. Additionally, unlike coverage that “write more tests to increase the coverage”, there is no easy and reasonable way to reduce complexity and coupling.

Size is an exception because it naturally have no preference to be higher or lower and it is the only vital signs that does not have a default coloring method. It is intended to stay in default vital signs set as a reference rather than a indicator.

The second reason is that productive metrics are less dynamic than

The last reason, as indicated in students’ responses, is because Software ICU is used by some students to improve their team’s process by tracking members’ activities. As mentioned by a student, member-level Telemetry analyses provide a quantitive way to identify who is falling behind in terms of effort output, thus team members can be more self-critical by comparing their individual data to the groups. Students do make use of these process vital signs to better organize team development. And these vital signs offer a way to motivate students to work hard.

However, it is concerned that the DevTime and Commit are so popular that it may also imply high risk of measurement dysfunctions that affect user’s behaviors. As noted by Austin in his Measuring and Managing Performance in Organizations[?], measurement dysfunctions defining characteristic is that the actions leading to it fulfill the letter but not the spirit of the stated intentions. At least one student actually experienced this negative effect. He explicitly point out that the quantitative measurement of their activities lead to a competition of stats within the group. More students have possibly been affected as well because as indicated in that student’s answers, his team share the similar opinion of the “stats competition”.

## **5.7 Feasibility in a professional software development context**

Responses of the questionnaire show that most students thought it was at least somewhat feasible to use Hackystat and Software ICU as a professional developer. Most comments indicate

that the system would be helpful in professional settings. There are some arguments about the potential bias of analysis data of Hackystat that the statistical data did not accurate enough to exclusively determine a project's health state in Software ICU.

## **Chapter 6**

# **Contribution and Future Directions**

## Appendix A

# 2008 Classroom Evaluation Questionnaire of Hackystat

### *Hackystat Evaluation*

*Hackystat is a long term research project concerned with improving the effectiveness and efficiency of software engineering metrics collection and analysis. Since 2003, we have periodically conducted a survey of students in ICS software engineering classes to assess the current strengths and weaknesses of the system.*

*To preserve anonymity, while also ensuring that only ICS students respond and respond only once, we ask you to provide the "secret code" that you randomly selected in class. To enable credit for completing this evaluation, only the graduate student researcher on this project (Shaoxuan Zhang) will know which code corresponds to you. He will provide a list of names who should be awarded credit to the class instructor without identifying individual responses. You can also contact Shaoxuan if you want your data deleted from analysis after you've submitted it.*

*If you want to go back and change your responses, simply fill out the entire form again. We will discard all but the most recently submitted entry for a given code.*

*This survey contains 17 questions and we expect that you will need about 10 minutes to complete it.*

*Thank you very much for your help! We take your views very seriously: prior responses to this survey have led to far-reaching changes in Hackystat.*

*Before filling out this questionnaire, you might want to take a look at the following image for the Software ICU to refresh your memory:*

<http://csdl.ics.hawaii.edu/~johnson/portfolio.gif>

*\* Required*

1. Installing the Eclipse IDE sensor was: \*
- Very Easy
- Easy

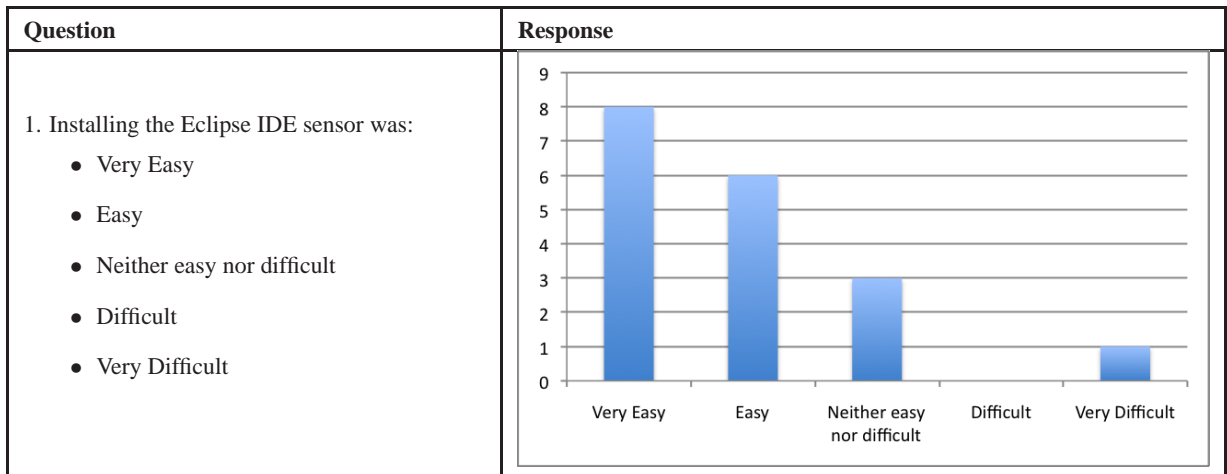
- *Neither easy nor difficult*
  - *Difficult*
  - *Very Difficult*
2. *Installing the Ant sensors (JUnit, SCLC, Emma, etc.) was: \**
    - *Very Easy*
    - *Easy*
    - *Neither easy nor difficult*
    - *Difficult*
    - *Very Difficult*
  3. *Please provide any feedback you can on the problems you experienced during sensor installation and server configuration, as well as any suggestions you have to make this easier in future.*
  4. *The amount of overhead required to collect Hackystat data (after successful installation and configuration of sensors) was: \**
    - *Very Low*
    - *Low*
    - *Neither low nor high*
    - *High*
    - *Very High*
  5. *The amount of overhead required to run Hackystat analyses was: \**
    - *Very Low*
    - *Low*
    - *Neither low nor high*
    - *High*
    - *Very High*
  6. *Please provide any feedback you can on Hackystat overhead, as well as any suggestions you have to reduce the overhead in future.*
  7. *Did you encounter any problems while collecting data? Was there any kind of data that you failed to collect? If yes, please explain.*
  8. *How did you feel about sharing your software development data with other members of the class? \**
  9. *How frequently did you use the telemetry page? \**
    - *Every day or more*
    - *2-3 times a week*
    - *Once a week*
    - *Less than once a week*
    - *Never*
  10. *If you used the Telemetry page, what were you trying to find out?*
  11. *How frequently did you use the Software ICU? \**

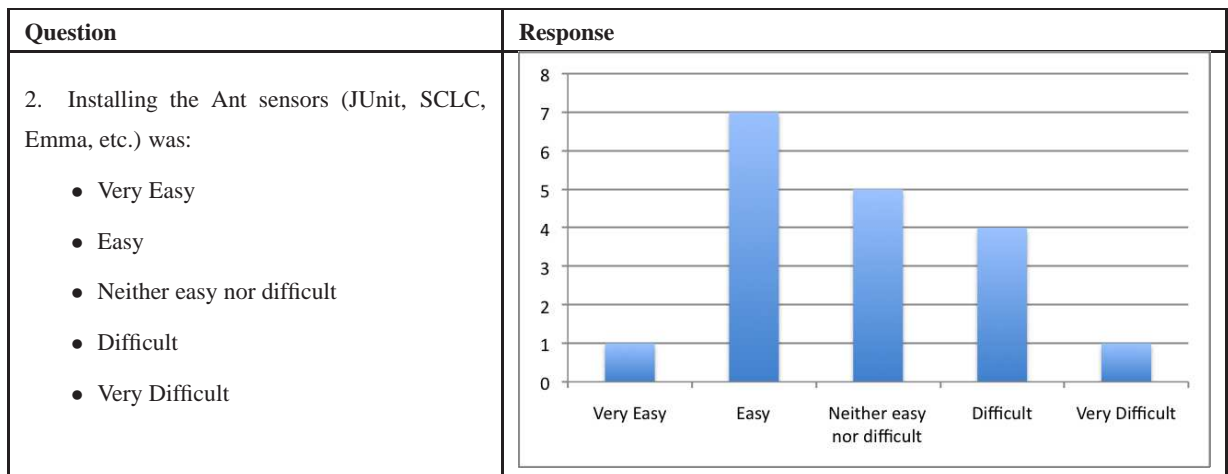
- *Every day or more*
  - *2-3 times a week*
  - *Once a week*
  - *Less than once a week*
  - *Never*
12. *If you used the Software ICU, please check the vital signs that were useful to you.*  
\*
- *Coverage*
  - *Complexity*
  - *Coupling*
  - *Churn*
  - *Size*
  - *DevTime*
  - *Commit*
  - *Build*
  - *Test*
  - *None of the above*
13. *Did you feel the Software ICU colors accurately reflected the "health" of your project? If not, why not? \**
14. *Were you able to use the Software ICU to improve your software's quality and/or your team's process? If so, in what ways? If not, why not? \**
15. *Please provide any other feedback you would like regarding Telemetry and the Software ICU, as well as any suggestions you have on how we can improve the system.*
16. *If I was a professional software developer, using Hackystat at my job would be: \**
- *Very feasible*
  - *Somewhat feasible*
  - *Neither feasible nor infeasible*
  - *Somewhat infeasible*
  - *Very infeasible*
17. *Please provide any other feedback you can on the feasibility of Hackystat in a professional setting, as well as any suggestions you have on how its feasibility could be improved.*

## Appendix B

# Results form 2008 Classroom Evaluation Questionnaire of Hackystat

This section presents the responses from the respondents to each of the questions. For the "short answer" questions, I corrected misspellings and minor grammatical errors to improve readability.





3. Please provide any feedback you can on the problems you experienced during sensor installation and server conguration, as well as any suggestions you have to make this easier in future.

- I could not figure out what step makes a .hackystat directory. My .hackystat directory automatically generated in my Documents and Settings directory which has a blank space in directory name. I am still not sure how to move this folder to other. The installation of all sensors was pretty well described at the project homepage and there was no problems I have met during the installation.
- Both the installation and sending sensor data was easy. However, tracking down whenever there is a problem with the sensor is not so easy. A troubleshooting page in the near future?
- Installing the sensors was pretty straightforward. I didn't have any problems.
- Case sensitivity was one problem between user and Hackystat, but it was fixed.  
If it is possible to have a .EXE that will automatically create environment variables and also install files into a local directory will be awesome.
- I did have one small hang up when installing the Ant sensors: If I remember correctly I was getting a NoClassDefinition error whenever a sensor ran. I was running java 1.5. I fixed it by downloading the jaxb libraries since the errors were referring to that. It could be not related to jaxb at all, but it worked after that. Otherwise, I had no problems whatsoever installing the sensors.
- Everything went smooth with the instructions given and the verification after each step.



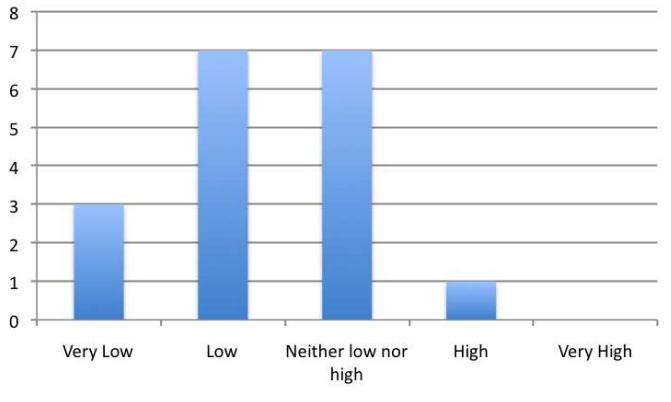
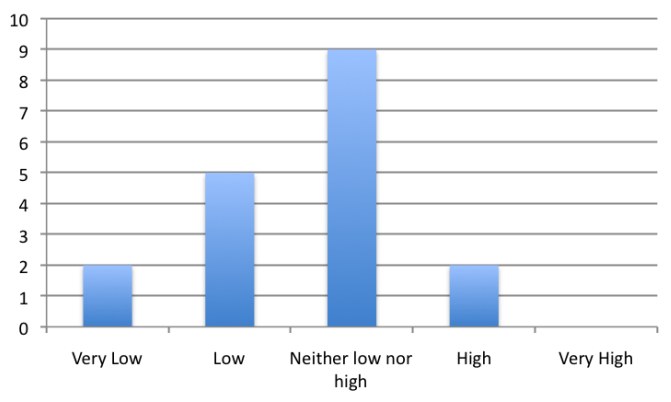
- Personally I didn't run into any problems but some of the other students did. The sensors aren't difficult to install per se, but there are a lot of steps involved and it's easy to get lost while installing them. Maybe an automated installer can be created that searches for the Ant tools (maybe the user can provide a search directory) and will configure and install the sensors for the user.
- What made it hard was that all the instructions were not in one page. I had to go from one page to another and then to another. There should be instructions from STEP 1 to the end and provide proper links to the step by step process.
- First of all, the manual is too long. I do like your goal to analyze the software project, but if it wasn't required by this class, maybe I wouldn't think I want to use it, because it looks too complicated.

Also there are too many things that we need to download and install. If you want to encourage people to use this more, maybe you should provide a package of all the tools somehow.

For example, before it took a long time to install Apache, MySQL, PHP, and Perl, but now somebody offers a package called XAMPP, which is a combination of all of those, and entire installation finishes in 3 minutes. Something like that should be given.

- There is a lot of documentation in a lot of different places. It was confusing trying to figure out what to read in what order, and whether or not it was relevant to me.
- Some the installation instructions could benefit from "write once, use many times" as they're repeated, which causes some people to start glossing over the instructions and then there's a couple that are slightly different and people (like me) won't notice the difference.
- The walkthrough was great, which made the installation easy.
- The only problem I had was the installation of the ant sensor. I mean configuring it on Eclipse was easy especially when I try to run Emma, JUnit, FindBugs and all that from Eclipse it is sending stuff to Hackystat but when I checked my software ICU I didn't have any data on Build (all it says was N/A). And little did I know that when you run the ant sensors on Eclipse it only registers all the data to Hackystat JUnit, Emma, Checkstyle and such except BUILD. And I was told that running the BUILD on the command line works but not on Eclipse. So I tried that and YES that works. So is there a way to make it work on Eclipse when you run all the ant sensors and it sends all the data to Hackystat including the BUILD data?

- When we ran the svn sensor, the build would fail if there are any commits from members not identified in our local Usermap.xml. Instead of looking for all commit records from all users within 24 hours, perhaps it could filter out and only look for records inside our UserMap.xml.
- The installation documentation must be read carefully. It may be easier to create a hackystat.build.xml with all the build targets, then import that file into each \*.build.xml and call the sensor from the tasks.
- The most challenging sensor to get up and running was the SVN sensor. Other than that, the others seemed fairly easy to install.

Question	Response												
<p>4. The amount of overhead required to collect Hackystat data (after successful installation and configuration of sensors) was: *</p> <ul style="list-style-type: none"> <li>• Very Low</li> <li>• Low</li> <li>• Neither low nor high</li> <li>• High</li> <li>• Very High</li> </ul>	 <table border="1"> <thead> <tr> <th>Response Category</th> <th>Count</th> </tr> </thead> <tbody> <tr> <td>Very Low</td> <td>3</td> </tr> <tr> <td>Low</td> <td>7</td> </tr> <tr> <td>Neither low nor high</td> <td>7</td> </tr> <tr> <td>High</td> <td>1</td> </tr> <tr> <td>Very High</td> <td>0</td> </tr> </tbody> </table>	Response Category	Count	Very Low	3	Low	7	Neither low nor high	7	High	1	Very High	0
Response Category	Count												
Very Low	3												
Low	7												
Neither low nor high	7												
High	1												
Very High	0												
<p>5. The amount of overhead required to run Hackystat analyses was: *</p> <ul style="list-style-type: none"> <li>• Very Low</li> <li>• Low</li> <li>• Neither low nor high</li> <li>• High</li> <li>• Very High</li> </ul>	 <table border="1"> <thead> <tr> <th>Response Category</th> <th>Count</th> </tr> </thead> <tbody> <tr> <td>Very Low</td> <td>2</td> </tr> <tr> <td>Low</td> <td>5</td> </tr> <tr> <td>Neither low nor high</td> <td>9</td> </tr> <tr> <td>High</td> <td>2</td> </tr> <tr> <td>Very High</td> <td>0</td> </tr> </tbody> </table>	Response Category	Count	Very Low	2	Low	5	Neither low nor high	9	High	2	Very High	0
Response Category	Count												
Very Low	2												
Low	5												
Neither low nor high	9												
High	2												
Very High	0												

6. Please provide any feedback you can on Hackystat overhead, as well as any suggestions you have to reduce the overhead in future.

- Since the verify command runs all the tests, I'd think that it should send data for all tests run. Rather, in the portfolio analysis, the Unit Test portion only retrieved data for any JUnit builds

that were run. It doesn't really make sense why we'd have to run it separately when verify does it anyway.

- If I am correct, overhead - the processing time required by a device prior to the execution of a command. Then it all depends on what computer the user is using, I am using a single-core processor laptop it did not take long.
- Since Dr. Johnson provided us with Ant sensor examples, it was quite easy to set up everything to send data to the sensorbase. I did the hackystat tutorial and everything worked fine. However, I missed the part about creating a usermap.xml file for the svn sensors through Ant. That confused me a bit later on but I figured it out.

What made getting data quite easy as well was having Hudson installed on a dedicated continuous integration server. Daily builds would auto-send data to Hackystat and this made it super easy to get daily info.

- The sensors ran automatically and it was fast with sending the data.
- Maybe there can be a link on <http://dasha.ics.hawaii.edu> to both the Hudson and Hackystat server, that way we don't have to memorize the port numbers. Also, allowing us to create an account and password would go a long way towards usability. I had to put the Hackystat login information in a text file because I can't remember a randomly-generated string for the password.
- Sending sensor data was often quite slow. Generating reports in the web application was sometimes also slow – the page wouldn't load until you refreshed it.
- The overhead to collect data was generally small, however long enough that would generally run multiple (DOS) terminals so that I could continue working while it was sending data. Analysis was no overhead since that was just pulling up a browser page.
- When sending hackystat data, it was fairly quick on my computer, MacBook Pro. Tho, there were some students I saw which had a LONG wait time on the same laptop.
- I love Hackystat! It is a very great tool especially for a developer like me.
- Since Ant takes care of running Hackystat sensors, this made it very easy to accomplish.

7. Did you encounter any problems while collecting data? Was there any kind of data that you failed to collect? If yes, please explain.

- I had a problem with sending commit data to hackystat when I worked on a group project. That was because I did not update my sensors to newer version.
- At first during the implementation of DueDates 2.0, it was not collecting commit data from my account. It was due to the account on hackystat, it included the @gmail.com part of my gmail account. So it was not matching up with each other, the hackystat account and my gmail account.
- Running an analyses on my machine was slow, it would take over 3 minutes to run a build. I am not sure why it took so long to send the build data so I can't make a suggestion.
- Only JUnit data as mentioned previously.
- Case sensitivity was an issue at first, but it was corrected so I did not get problems after that. Hudson did not send to Hackystat number of commits, but that was fixed after a little modification with build.xml file.
- I was lucky. I rarely had any problems collecting data during all the time I worked with Hackystat. The one time something got screwed up was with my development time for one day. It said 0 when I checked and I had put in a bunch of time that day so it should have said otherwise.  
  
I don't remember exactly but, that night I believe had worked in eclipse till after 12 at night, so it went to the next day before I closed the program. That could possibly be a reason for the missing data initially. The next day I just cleared the cache and it was all fine.
- There was a small issue when I first started collecting data, but it was quickly corrected when checking the xml files.
- Personally I ran into no problems collecting data.
- Sometimes it didn't collect build data for some reason.
- Occasional problems with SVN collection, I think, was a bit hard to tell.
- Everything was great except collecting data for my BUILD (please refer to above statement for more detailed problem regarding this). Thank you.
- I did with commit records but it was my fault. I wish subversion with Google Project Hosting would be more strict. I was able to check out the project with or without the "@gmail.com" suffix (i.e. "test" and "test@gmail.com"). Thus making me two different authors.


- Yes, the build data. I needed to set more environmental variables.
- For some unknown reason, my user name picked up the @gmail.com, so both my user name with and without @gmail.com needed to be added to the projects.

8. How did you feel about sharing your software development data with other members of the class?

- We could see how other groups were doing by sharing our software development data with other people. We also could find out what kinds of problems with our project by comparing graphs with other groups and this helped a lot.
- I was not offended if it was low, and was quite intrigued with others data.
- I did not have a problem with sharing data with other people in class. I thought it was needed tool to keep tabs on everyone to assure they're doing their fair share.
- It felt good if your data was better than others. And if it wasn't, then you felt bad.
- Did not really like it because it is showing my programming habits, like starting on a project on the last couple of days.
- I felt alright about sharing my data with the class. It was interesting for me to see how other people worked on stuff. Some were consistent and others were not. Some people spend a lot of time working on stuff yet do not commit as much as others that work half the time. I think its good to see this data.
- I am okay with sharing my data.
- I didn't think it was a particularly good idea because it then forces group members to become competitive with each other, especially if one person is able to put in more time than all the others. Also, the data doesn't reflect the amount of work put in, maybe someone spent 5 hours doing research and only 1 hour programming, but the sensor data will only show 1 hour of development time and a minor code commit, versus someone who, say, just changes around the package structure for 3 hours and has a huge commit amount.
- Actually hackystat (or hacky-stalk as what my teammates and I called it) caused a lot of arguments and trash talk. Some guys were more concerned about collecting stats on hackystat than actually finishing the project. Some members would start competing on who had more

commits or move development time. The project turned out to be more of a competition of stats, which wasn't healthy for the team at all.

- It will be obvious that who worked on the project, so it is nice in terms of grading students. At the same time I feel some pressure that I need to work on the development, so if team leader require everybody to work well, this is good.
- Didn't really care.
- I had no problem with this, and it encouraged me to be aware of my time management and coding style.
- It was good in a sense that they can help you with test cases and coverage.
- It was fun..because you can see how everyone is doing within your group.
- Before taking this class, I didn't think that there was a way to track software development process. After learning about software continuous integration and working in a larger group project, I have a better insight in sharing the development process. I feel that it is a must in every software development environment, big or small to be able to communicate frequently and effectively.
- I was nervous because certain individual of the class seemed able to put in ridiculous long hours. I was concerned my amount of time (which seemed reasonable) would make me look as though I'm not working as hard.
- Good, I can see how I and others rank with each other.
- I am fine with this. All group projects in all schools (e.g., Architecture) should be required to use such a system. This is great for facilitating fair evaluations of students who participate, and those who 'get the grade' by riding on the laurels, blood, sweat, and tears of others.



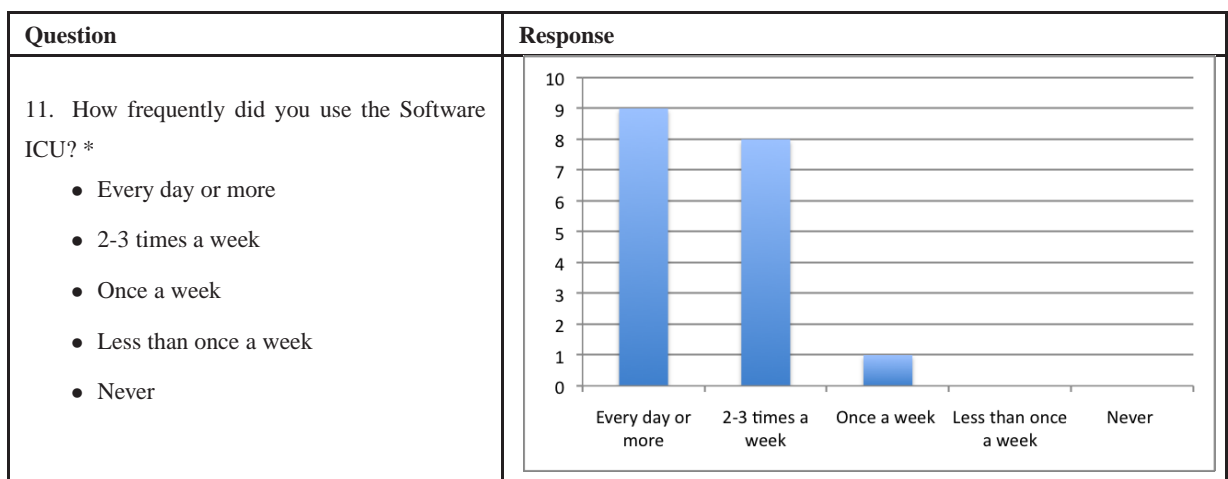
10. If you used the Telemetry page, what were you trying to find out?

- I tried to find out how was I doing for the project by looking hackystat data.
- Seeing how much time i spent on the development of the program, and also others in my group.
- When I used the telemetry page I was trying to find out if I was on par with other groups members in terms of development, build, and commit numbers.
- Whether or not, my sensors were reading, and the work output of my group members (especially on days we didn't meet together).
- If my development time was up to par with my team members.
- I usually used the telemetry page to evaluate how my team was working overall, and what my part was in that data. I also checked it to make sure everyones data was being sent.
- It helps me see how I measure up with my partners.
- Member dev time mostly, to compare the amount of development time I put in vs. my group members.
- It supposed to show us how healthy individuals are in the group. So if one person is slacking, the members need to tell him to step it up. It wasn't used that way in our group. One person really wanted a good grade for the class so he just used the telemetry to watch himself; making sure no one gets more builds/devTime/commits than him (yes he said "i need more dev time because i need an A"). I remember we had dinner as a group and one of our group members

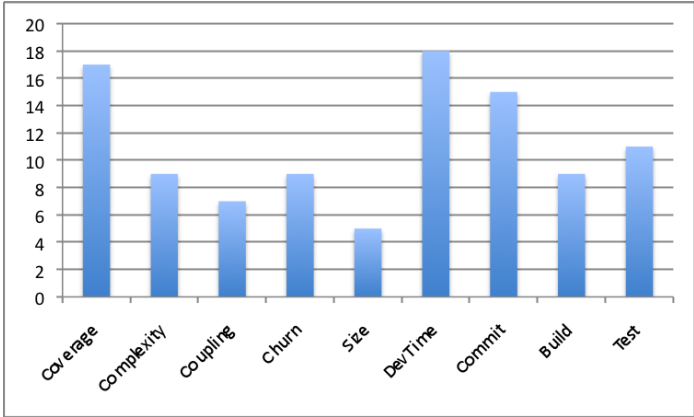
didn't go to dinner. another group members then said "oh if he ups his stats more than mine, tomorrow I'm gonna hack all day."

Sad, but true.

- member commit, member dev time
- Curious about trends in dev time, commits.
- Usually MemberDevTime, MemberBuilds, and MemberCommits. Basically just seeing how everyone was progressing.
- graphs, line trends of other group members
- My status and the status of our group and make sure everyone is doing their part.
- Mostly trends in individual performance, as well as overall project outlook.
- Basically if everyone was putting in the same amount of effort. Also it helped indicate if everyone is on track. If they have regular activity, then the chances of them on track is higher.
- Was the coverage, complexity and coupling getting bad?
- I tried to review each telemetry page daily to understand what I could do to improve the project health and focus efforts.





Question	Response																				
<p>12. If you used the Software ICU, please check the vital signs that were useful to you. *</p> <ul style="list-style-type: none"> <li>• Coverage</li> <li>• Complexity</li> <li>• Coupling</li> <li>• Churn</li> <li>• Size</li> <li>• DevTime</li> <li>• Commit</li> <li>• Build</li> <li>• Test</li> <li>• None of the above</li> </ul>	 <table border="1"> <thead> <tr> <th>Vital Sign</th> <th>Frequency</th> </tr> </thead> <tbody> <tr> <td>coverage</td> <td>17</td> </tr> <tr> <td>complexity</td> <td>9</td> </tr> <tr> <td>coupling</td> <td>7</td> </tr> <tr> <td>churn</td> <td>9</td> </tr> <tr> <td>size</td> <td>5</td> </tr> <tr> <td>DevTime</td> <td>18</td> </tr> <tr> <td>Commit</td> <td>15</td> </tr> <tr> <td>Build</td> <td>9</td> </tr> <tr> <td>Test</td> <td>11</td> </tr> </tbody> </table>	Vital Sign	Frequency	coverage	17	complexity	9	coupling	7	churn	9	size	5	DevTime	18	Commit	15	Build	9	Test	11
Vital Sign	Frequency																				
coverage	17																				
complexity	9																				
coupling	7																				
churn	9																				
size	5																				
DevTime	18																				
Commit	15																				
Build	9																				
Test	11																				

13. Did you feel the Software ICU colors accurately reflected the health of your project?  
If not, why not?

- I felt most of colors accurately reflected the health of the project. For the Coverage data, since we can write test cases just for increasing of the rates, we cannot assume that the project is in healthy condition even if the coverage data displayed in green color. However, I think this is not a problem of hackystat.
- Yes
- The only issue I had with the ICU colors was with the coupling. In both versions of DueDates we had to add extra classes at the last minute which would cause the coupling ICU to turn red. I am not sure how to address that because the coupling does need tracking.
- Not really, I don't think having a high churn amount is necessarily bad. Of course, it's a case-by-case thing. For my group, it wasn't about not committing frequently; we were just rehashing code because something just didn't work.
- Yes, reflected accurately on the health of the project. Showed how much coverage we had.
- I feel that the Software ICU did accurately reflect the health of my projects. For Due Dates 2.0, which was a longer project, the data was getting increasingly more meaningful as the

trends were over a larger period of time. It is good to look at things like devtime, commits, coupling, and coverage to see the color and the past trend because i think they really say something about the current state of the project.

To make it simpler, whenever I knew our project wasn't doing good and people weren't working regularly, the software ICU would have lots of reds and yellows. When I knew the project was doing better and people were working regularly, there were greens. It makes sense.

- The ICU was accurate with our project because it showed drastic spikes in all signs. This reflects our project in poor health.
- Not particularly because a project's health cannot easily be determined by just measuring numbers alone. For example, it's easy to increase coverage, but if a class has nothing but getters and setters and a toString method, does it really need to be tested? Of course not, but someone might feel compelled to do it in order to increase coverage and get a better health, but it's just a waste of time in my opinion. Also, DevTime is only measured from Eclipse but that doesn't measure things such as someone reading a book or looking up websites for information. It only measures active development in one program, forcing people to only use whatever IDE's Hackystat supports. The figures for complexity and coupling are hard to evaluate too. We want complexity to be low but sometimes it's unavoidable for it to be high, and should Hackystat show an absolute cut-off point where the complexity must be below a certain point for the project to be considered acceptable? Coupling is another one that falls under this category, if your program relies on a lot of outside libraries, can someone really determine an absolute value that the project's coupling must be under?
- Yes.
- maybe
- Coverage: perhaps too sensitive to drops/bounces in coverage. Churn: while you're working on a project, churn is going to vary, sometimes a lot. The trend colors were not helpful.
- Yes, I felt it was a relatively healthy project, and this generally showed, in the end. In the first half the colors reflected not as health of a project, which I'd agree as well. I'm not sure rising coupling was entirely a bad sign as things went along and functionality was added, as it was a slow steady rise.
- Sometimes. Hard to determine what will fall into green, red, or yellow.

- Yes definitely.
- It somewhat reflected the quality of our project. Maybe in some dark corner something is not thoroughly being depicted through the colors. Perhaps a suggestion is to use different color hues.
- Yes it was pretty accurately reflected.
- No, since I did not correctly configure the sensors.
- This is subjective... Usually the colors were spot on, however, they are quick to turn one way or the other depending on events that are being managed by the team (e.g., large code churns due to removal of unused code/imported code, etc.).

14. Were you able to use the Software ICU to improve your software's quality and/or your team's process? If so, in what ways? If not, why not?

- We can check how other members are doing for the project through the Software ICU and this helps a lot especially when we are working on the team project.
- Yes, for tracking if members were working on their tasks. Also how complex the program is increasing or decreasing.
- In my opinion, it is not clear if the ICU improved our system. Because other tools such as junit, findbugs, and pmd was easier to use to improve the application.
- If anything, keeping an eye on coverage helped us look out for what was being tested and what wasn't. Yes, showed how much coverage we had, and improve on that.
- I think for sure the Software ICU improves team process. More than just keeping people "in check" when grades are at stake, it provides an accurate way to assess what's being done and by whom. Our team got a lot out of checking up on the software ICU and assessing our team process. It seemed to get better over time.

As far as the software's quality, I think the Software ICU could be very useful in improving this. If my project for instance was in the red for complexity and coupling, and there were some code issues, I could see all this automatically through hackstat. Besides coverage stats though, my team did not really use the ICU to improve the software's quality.

- ICU was able to help us because it told us what needs to be focused or corrected.

- Personally, I only found Hudson useful because it's like running your code on someone else's computer to see if your environment is set up differently from a generic machine. I feel that the data for Hackystat is more something to look at out of curiosity rather than something to determine how well a project's status is because it's hard to base a project's health based on numbers alone and it might put unrealistic pressures on the team to make the project healthy for Hackystat when they can better spend their time developing instead.
- Yes.
- Yes, coverage tells me if we didn't write enough test cases.
- No. Coverage: already aware from Emma. DevTime, Commit, Build, Test: either team members did not look at the statistics, or they didn't care, because their habits did not change much. Others: not much we could do about the other statistics.
- Yes because able to manage our time and development fairly equally, and also notice spikes indicating bigger changes or problems.
- Yes, shows where we could improve as a group and improve as a programmer.
- Like in my case last time, I saw on Software ICU that I don't have a data on my BUILD. So because of that information I know what the problem is and it helped me to find a solution and figure everything out before it is too late.
- Our project ICU definitely described our lacking and late attempt to improve coverage. Due to the ICU, we were able to distinguish this fact quick and easy.
- The amount of activity helped us identify who was falling behind. Without offending our members by outrageously claiming their not working, we could tell by the sensors. Members can be more self-critical by looking at their individual data compared to the groups.
- Yes, by checking the coverage, complexity and coupling.
- Yes. By targeting coverage, dev time, coupling, and complexity, my team was able to improve all these into areas that were acceptable to us.

15. Please provide any other feedback you would like regarding Telemetry and the Software ICU, as well as any suggestions you have on how we can improve the system.

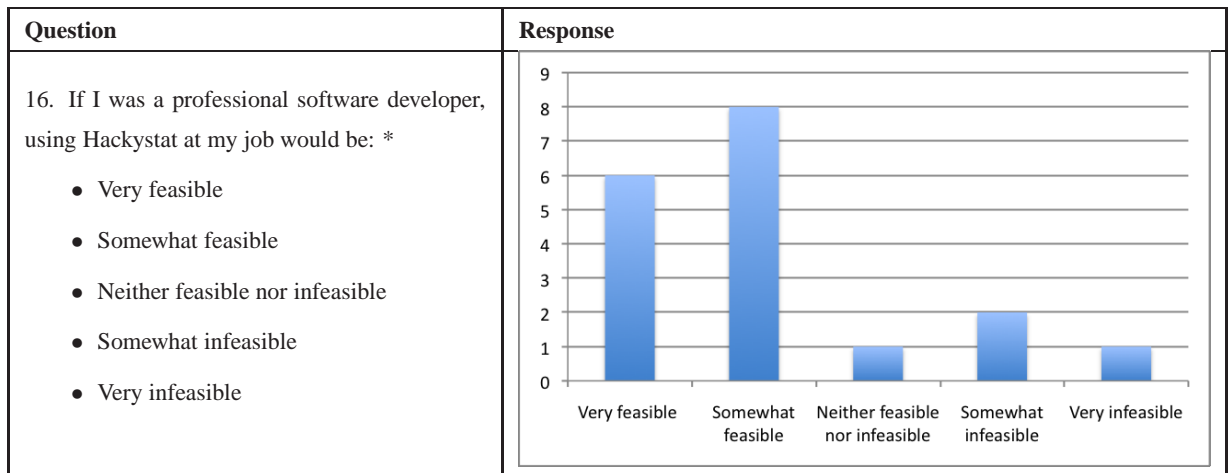
- I do not think the commits, builds, tests should be colored in because it all depends on how much the user does on the project. Is it possible to show line coverage instead of method coverage? The software ICU and telemetry was awesome tools in helping out with the project. It gave me visual stats on the project.
- What I think would be cool is to implement something to view the trend for each category in larger format but in the same style as the software ICU. I know this is shown on the telemetry page when you select it to show. However, I would be nice if there was some sort of rollover function that brought up a slightly larger window with a blown up overall trend. I can see how this isn't really needed but I would mostly likely check it a lot if it was there.

A minor thing that I noticed when using the Telemetry page was that when I selected a new statistic to view, the page would always jump back to the top and I'd have to scroll down each time. Its not really a biggie, but it makes navigating a bit slower when your going through all the project statistics.

- Consistent colors for each members can help.
- In addition to everything I mentioned above, it might help to somehow make the sensors configurable in some way, for example if two people are doing pair programming, there should be an option to set the sensors to send data for both people. Perhaps complexity can be measured somehow to only include methods that, say, start with get or set and toString. This way people aren't forced to write pointless test cases in order to increase coverage.
- Help page should be provided inside project browser. It should describe how to use it, what telemetry, what churn is, something like that.

Also your explanation should be simple so that people want to read it. If it is complicated and long explanation, nobody will read it.

- The different color bars and randomness might be fun and interesting, but I think having a bit more consistent scheme might be better. I would suggest if possible giving each developer a specific color that they always have during the project, either random, or chosen at the beginning.
- Does not capture development outside of Eclipse. For example, IMHO, MS Visual Studio is much better in the capacity as a web development IDE, which the dev time here was not recorded.



17. Please provide any other feedback you can on the feasibility of Hackstat in a professional setting, as well as any suggestions you have on how its feasibility could be improved.

- I think its good to have this in a professional environment, cause the employer or client can check on how the progress of the program is going. With out having to make so much visits or hovering over workers.
- Cannot think of any off of my head. The Software ICU is already great for us programming students.
- I think Hackstat is definitely feasible in a professional setting, as long as it is supported in some way. For instance, if a team of developers is working on a project and they are all for having Hackstat manage project stats, that would be great. If, however, your the only person on your team that wants to use it, then it would be hard to send data that would assess team process.

I could see project managers wanting to have Hackstat data to evaluate everyone's input into the project, as well as the health of the project. Hackstat, I think, is perfect for new open source projects if releases are made early and often. It could be essential to seeing the overall health of the project.

- Overall, I feel like Hackstat would be an interesting tool to gather data to look at for curiosity's sake from time to time, but it should not be used as a basis for determining a project's health or to determine something such as member contribution. The sensors can only gather information from a few sources and these readings cannot account for a person's full contributions to a project. As for determining a project's health, I do not believe the sensor readings

can provide an accurate measurement because the sensors can only measure numbers based on algorithms, but it takes a person to really determine how good the code is.

- When I start to use hackystat, I need to get password from you and then eclipse send my data to your server. Some developers might have concern that hackystat steal source code.
- I think it depends a lot on the culture of job setting. I'm not too sure, but I think I may try setting it up on my own job site, even if just for myself to see my own trends.
- It is a very useful tool to keep track the health of a project so I would say it is feasible to have it in a job.
- My only wish is that ICU's should have a feature to support pair programming. Possibly a feature to indicate to the system that two people may be working on the same problem on the same system, rather than two individual machines. You might want to call this "collaborative mode," or something along the lines of that. These settings of course should be turned on or off easily from the developer's IDE (Eclipse).
- I work in a one person shop, so it would be difficult to say how useful this would be. As a lone developer, many metrics I am very cognizant of, however, having such a system would allow me to view those statistics that I do not have a "gut" feeling for. It would be great for my boss to measure the amount of time I spend on a project however.

# Bibliography

- [1] Irina Diana Coman, Alberto Sillitti, and Giancarlo Succi. A case-study on using an automated in-process software engineering measurement and analysis system in an industrial environment. In *Proceedings of the 2009 IEEE 31st International Conference on Software Engineering*, pages 89–99. IEEE Computer Society, 2009.
- [2] Coupling (computer science). [http://en.wikipedia.org/wiki/Coupling\\_\(computer\\_science\)](http://en.wikipedia.org/wiki/Coupling_(computer_science))
- [3] Emanuele Danovaro, Tadas Remencius, Alberto Sillitti, and Giancarlo Succi. Pem: experience management tool for software companies. In *Companion to the 23rd ACM SIGPLAN conference on Object-oriented programming systems languages and applications*, pages 733–734. ACM, 2008.
- [4] Google chart api. <http://code.google.com/apis/chart/>.
- [5] Watts S. Humphrey. *A Discipline For Software Engineering*. Addison-Wesley, New York, 1995.
- [6] Watts S. Humphrey. *Introduction to the Teasm Software Process*. Addison-Wesley, New York, 2000.
- [7] Philip M. Johnson. Results from the 2003 classroom evaluation of Hackystat-UH. Technical Report CSDL-03-13, Department of Information and Computer Sciences, University of Hawaii, Honolulu, Hawaii 96822, December 2003.
- [8] Philip M. Johnson. Results from the 2006 classroom evaluation of Hackystat-UH. Technical Report CSDL-07-02, Department of Information and Computer Sciences, University of Hawaii, Honolulu, Hawaii 96822, December 2006.
- [9] Philip M. Johnson, Hongbing Kou, Joy M. Agustin, Christopher Chan, Carleton A. Moore, Jitender Miglani, Shenyan Zhen, and William E. Doane. Beyond the personal software pro-



cess: Metrics collection and analysis for the differently disciplined. In *Proceedings of the 2003 International Conference on Software Engineering*, Portland, Oregon, May 2003.

- [10] Thomas J. McCabe. A complexity measure. *IEEE Transactions on Software Engineering*, 2(4):308–320, December 1976.
- [11] Alberto Sillitti, Andrea Janes, Giancarlo Succi, and Tullio Vernazza. Collecting, integrating and analyzing software metrics and personal software process data. In *Proceedings of the 29th Conference on EUROMICRO*, page 336. IEEE Computer Society, 2003.